

# Algorithms and Parallel Computing



**POLITECNICO**  
MILANO 1863

Course 052496

Prof. Danilo Ardagna, Prof. Matteo Giovanni Rossi

Date: 10-02-2025

Last Name: .....

First Name: .....

Student ID: .....

Signature: .....

Exam duration: 2 hours and 30 minutes

Students can use a pen or a pencil for answering questions.

Students are NOT permitted to use books, course notes, calculators, mobile phones, and similar connected devices.

Students are NOT permitted to copy anyone else's answers, pass notes amongst themselves, or engage in other forms of misconduct at any time during the exam.

Writing on the cheat sheet is NOT allowed.

Exercise 1: \_\_\_\_\_ Exercise 2: \_\_\_\_\_ Exercise 3: \_\_\_\_\_

## Exercise 1 (up to 13 points)

13 Points if you answer questions 1–4+6

10 Points if you answer questions 1–4+5

We want to realize a system for the acquisition of sequences of data (time series) of different kinds—for example, sequences of temperature readings from a sensor. We want the system to be general enough to handle different kinds of data, and also different strategies for the handling of the data (for example, for certain series it is crucial to handle the readings in strictly sequential order, while for other series data might be read out of order).

Answer the following questions.

1. Define a class **Reading**, which captures the general notion of reading.  
A reading is characterized by a timestamp and by the value that has been read. The timestamp is a value that corresponds to the number of milliseconds elapsed from the first of January 1970 until the time of the reading (i.e., 0 is the midnight of January 1st, 1970, 1 is one millisecond after that, and so on). The reading value is a real number.  
Once created, a reading cannot be modified.  
Class **Reading** has a single constructor, which initializes the timestamp and the value.  
Users of the class must also be able to retrieve the timestamp and the value of the reading.  
A reading can be valid or not. The class provides a function that returns **true** if the reading is valid, **false** otherwise. By default, a reading is valid (though for some kinds of reading the value must obey certain constraints, see the next 2 questions).
2. Define a class **TemperatureReading** that captures temperature readings (measured in Celsius degrees). To be valid, a temperature reading must be greater than or equal to -273 degrees.
3. Define a class **StockVariationReading** that captures readings concerning the variations (in percentage) of the value of a stock. Notice that a variation can be negative. We represent variations (in percentage) as real values between -1 and 1 (i.e., the value of a stock cannot increase or decrease by more than 100%).
4. Define a class **ReadingCollection**, which represents a collection of readings. Depending on the readings stored in the collection, the nature of the collection can vary. The only feature in common to all collections is that they offer an operation, **average**, which returns the average of the values of the readings stored in the collection.

Answer one (**and only one**) of the following two questions. Notice that the maximum total points of the exercise are different if you develop point 5 rather than point 6. In each case, in developing your solution, you should take into consideration the efficiency of operation **addReading**.

5. Define a class, **TemperatureCollection**, which is a collection of temperature readings. Note that readings could be shared among collections. Also, in the collection there must be a single reading per timestamp. The class provides an operation, **addReading**, which adds a new reading to the collection, and returns **true** if the addition is successful, **false** otherwise; temperatures **must** be added in order, according to their timestamps, otherwise the addition is not successful. The class also provides a method, **lastNvalues**, that takes as input a positive integer **n**, and returns the last **n** elements added to the collection (it returns all the elements in the collection if it contains less than **n** readings). Declare all elements of the class, and define operations **addReading** and **lastNvalues**.
6. Define a class, **StockVariationCollection**, which is a collection of stock variations. Note that readings can be shared among collections. Also, in the collection there must be a single reading per timestamp. The class provides an operation, **addReading**, which adds a new reading to the collection, and returns **true** if the addition is successful, **false** otherwise; stock variations can be added out of order. The class also provides a method, **maximumSubsequence**, that returns the sum of the readings of the subsequence of the collection whose sum of the values is maximum. For example, if the readings are (in order of timestamp, which is not reported here for simplicity)  $[0.04, 0.1, -0.2, 0.05, -0.02, 0.08, 0.07, 0.08, -0.02, -0.03]$ , the sum of the maximum sequence is 0.26, which is the total of the subsequence  $[0.05, -0.02, 0.08, 0.07, 0.08]$ . If the collection is empty, **maximumSubsequence** returns 0. Declare all elements of the class, and define operations **addReading** and **maximumSubsequence**. What is the worst case complexity of your implementation of operation **maximumSubsequence**?

**NB:** The implementation of operation **maximumSubsequence** does not need to be optimal from the point of view of its complexity.

## Solution 1

1. Declaration of class **Reading**:

```

1  class Reading {
2  public:
3      Reading(double v, unsigned ts) : value(v), timestamp(ts) {}
4      virtual bool is_valid() const { return true; }
5
6      double get_value() const { return value; }
7      unsigned get_timestamp() const { return timestamp; }
8
9  protected:
10     const double value;
11     const unsigned timestamp;
12
13 };

```

Since a reading, after being created, cannot be modified, the fields of the class are all **const**. Hence, they must be both initialized at construction time. They are declared with visibility **protected** to allow derived classes to directly access them. Notice that an object of type **Reading** behaves in a similar manner as a base type (such as **int** or **char**): it can be exchanged (copied, passed as parameter, etc.) freely, as if it were a simple value, because it is immutable.

The getter methods, **get\_value** and **get\_timestamp** have been introduced to allow classes other than derived ones to read the values of the fields. The behavior of method **is\_valid** depends on the nature of the reading, so its definition will be overridden by derived classes, hence it is declared **virtual**. Without further information about its nature, the reading is considered to be valid, hence the definition of **is\_valid** in the base class of the hierarchy of readings simply returns **true**.

2. Declaration of class **TemperatureReading**:

```

1  class TemperatureReading : public Reading {
2  public:
3      TemperatureReading(double v, unsigned ts) : Reading(v, ts) {}
4      bool is_valid() const override { return value > -273; }
5  };

```

A `TemperatureReading` is a special case of `Reading`, hence it inherits from the latter. The only difference between a `Reading` and a `TemperatureReading` is that a temperature must be greater than  $-273$  to be valid, hence method `is_valid` overrides the one defined in the base class and is defined accordingly.

3. Declaration of class `StockVariationReading`:

```
1 class StockVariationReading : public Reading {
2 public:
3     StockVariationReading(double v, unsigned ts) : Reading(v, ts) {}
4     bool is_valid() const override { return -1 <= value && value <= 1; }
5 };
```

Similarly to `TemperatureReading`, `StockVariationReading` inherits from `Reading` and overrides the definition of `is_valid`; in this case, a `StockVariationReading` is valid when it belongs to interval  $[-1, 1]$ .

4. Declaration of class `ReadingCollection`:

```
1 class ReadingCollection {
2 public:
3     virtual double average() const = 0;
4 };
```

We do not know anything about a `ReadingCollection` except that it provides a method, `average`. However, since we do not have any information about the intended behavior of a collection, we declare `average` as a *pure virtual* function.

Notice that `ReadingCollection` is, essentially, what is known as an *interface*: it provides the signature of operations (in this case, `average`), but nothing else. Any class implementing `ReadingCollection` will have to provide not only the definition of method `average`, but also the definition of the attributes.

5. Declaration of class `TemperatureCollection`:

```
1 class TemperatureCollection : public ReadingCollection {
2 private:
3     std::vector<std::shared_ptr<TemperatureReading>> readings;
4 public:
5     bool addReading(const TemperatureReading &tr) {
6         if (readings.empty() || readings.back()->get_timestamp() < tr.get_timestamp()) {
7             readings.push_back(std::make_shared<TemperatureReading>(tr));
8             return true;
9         } else
10            return false;
11    };
12
13    std::vector<std::shared_ptr<TemperatureReading>> lastNvalues(unsigned n) const {
14        std::vector<std::shared_ptr<TemperatureReading>> res;
15        if (readings.size() <= n)
16            res = readings;
17        else {
18            for (unsigned i = readings.size() - n; i < readings.size(); i++) {
19                res.push_back(readings[i]);
20            }
21        }
22        return res;
23    }
24
25    double average() const override;
26
27 };
```

Class `TemperatureCollection` implements interface `ReadingCollection`. As such, it defines how the collection is actually stored in memory. In the case of a collection of temperature readings, we know that, by construction, data must be inserted in the collection in strictly sequential order. Hence, we need a data structure for which

appending element at the end of the sequence is efficient. For this reason, a `TemperatureCollection` relies on a `vector` to store the readings. Indeed, the `push_back` operation is, on average, very efficient for `vectors`—except when a reallocation of the sequence is necessary, which, however typically occurs infrequently. In addition, since we know that readings are shared among instances of the class, to avoid wasting memory we store them on the *heap*, and we use *shared pointers* (rather than regular pointers) so that we do not need to concern ourselves with the details of the handling of the allocated memory (in particular, its deallocation). Notice that since readings are immutable by construction, the use of shared pointers to access the data does not pose any concerns regarding side effects and unwanted uncontrolled changes, which are simply impossible.

The `addReading` method makes sure that readings are inserted in ascending order of timestamp: if the timestamp of the element to be added is not greater than the biggest timestamp appearing in the collection (which, by construction, is the one of the last reading of the collection), the element is not added (and `false` is returned). The argument of function `addReading` is declared as a reference to `const` for clarity's sake; however, the object is by construction immutable, so it cannot be modified anyway; in addition, the object is a small one (it includes only 2 attributes, both of simple types) so, according to the guidelines seen in class, it could have been passed also by value, rather than by reference.

Function `lastNvalues` simply copies the pointers to the last N elements of the vector in a new vector, and returns the latter.

#### 6. Declaration of class `StockVariationCollection`:

```

1  class StockVariationCollection : public ReadingCollection {
2  private:
3      std::map<unsigned, std::shared_ptr<StockVariationReading>> readings;
4  public:
5      bool addReading(const StockVariationReading &svr) {
6          auto res = readings.insert(std::pair<unsigned, std::shared_ptr<StockVariationReading>>(svr.
              get_timestamp(), std::make_shared<StockVariationReading>(svr)));
7          return res.second;
8      };
9
10     // This is a very simple implementation of the function.
11     // It performs 2 nested loops, and in this way it checks, one by one, all
12     // subsequences of the collection (if the sequence is made of elements [e1, e2, e3], for example,
13     // it checks [e1], then [e1, e2], then [e1, e2, e3], then [e2], then [e2, e3], then [e3]
14     // for each sequence it computes the sum of its elements, and it keeps track of the
15     // maximum sum.
16     double maximumSubsequence() const{
17         if (readings.empty())
18             return 0;
19         double max = readings.cbegin()->second->get_value();
20         for (auto it = readings.cbegin(); it != readings.end(); it++) {
21             double sum = it->second->get_value();
22             if (sum > max)
23                 max = sum;
24
25             auto it2 = it;
26             for (it2++; it2 != readings.end(); it2++) {
27                 sum += it2->second->get_value();
28                 if (sum > max)
29                     max = sum;
30             }
31         }
32         return max;
33     }
34
35     double average() const override;
36
37 };

```

In the case of `StockVariationReadingCollection`, since we know that elements can be inserted out of order of timestamp, a `vector` is no longer the best solution. In fact, the ordering of the elements according to their timestamps is a crucial feature of the collection (for example, to compute the maximum subsequence); in principle, one could store them out of order and then sort them whenever needed, but a better solution is just to store them in order, as long as the storing is efficient. Since inserting an element in order in a sequence typically requires shifting  $O(n)$  elements of the sequence (with  $n$  the number of stored elements), using a `vector` is not the best solution. An ordered `map` (where the key is simply the timestamp of the reading), instead, fulfills all our needs. As in the case of a `TemperatureCollection`, we store the readings on the heap and use shared pointers to ease memory management.

An ordered map already guarantees that stored keys are all distinct, so method `addReading` simply needs to return the value returned by `insert` (more precisely, the second element of the pair returned by `insert`). The argument of function `addReading` is declared as a reference to `const` for clarity's sake; however, the object is by construction immutable, so it cannot be modified anyway; in addition, the object is a small one (it includes only 2 attributes, both of simple types) so, according to the guidelines seen in class, it could have been passed also by value, rather than by reference.

As mentioned in the comments before method `maximumSubsequence`, the function essentially goes through the subsequences of the collection one by one (notice that, since we are using a map indexed on the timestamps, the iterators return the elements of the collection in order of timestamp).

The complexity of the implementation of function `maximumSubsequence` is  $O(n^2)$ , with  $n$  number of elements in the collection. Indeed, the  $i$ -th execution of the internal for loop (with  $i = 1, 2, \dots, n$ ) performs  $n - i + 1$  iterations, so the total number of subsequences checked is  $\sum_{i=1}^n i = \frac{n(n+1)}{2}$ , which is  $O(n^2)$  (notice that the check of each subsequence, which is performed by lines 27-29, takes constant time).

Notice that function `maximumSubsequence` can be implemented in a much more efficient way. Indeed, the following implementation has, as the name suggests, linear complexity (i.e., it has complexity  $O(n)$ ):

```

1  // This alternative implementation performs a single pass, so it has linear complexity.
2  // More precisely, it keeps track of the current sum of the elements.
3  // The moment when the sum becomes negative it means that, if we continue adding
4  // values to the sum, in any case the single value of the next element will
5  // be greater than the addition of the current sum with the next element, hence
6  // we need to make sure that we "restart" computing the sum; for this reason variable sum
7  // is reset to zero, because when the next element is added to sum, the new value
8  // of sum will simply be the value of the next element.
9  double maximumSubsequence_linear() const{
10     if (readings.empty())
11         return 0;
12     auto it = readings.cbegin();
13     double sum = it->second->get_value();
14     double max = sum;
15     for (it++; it != readings.end(); it++) {
16         sum += it->second->get_value();
17         if (sum > max)
18             max = sum;
19         if (sum < 0)
20             sum = 0;
21     }
22     return max;
23 }
```

It is clear that `maximumSubsequence_linear` has linear complexity, since it performs a single loop through the elements of the collection.

## Exercise 2 (13 points)

You are required to develop a parallel program that determines whether a matrix of integers is diagonal. The input matrix is stored in a file (whose name is provided as a single command-line argument) and is accessible only by the process of rank 0. Ensure that you manage command-line arguments appropriately, raising errors as necessary. We want to exploit a parallel architecture to realize the program:

1. Describe in your own words what strategy you want to use to parallelize the computation (i.e., how you distribute data across processes, what computation each process performs, how you generate the results).
2. Write suitable MPI code that realizes the program above.

**Note 1:** you may assume that the input is a square matrix (no input size checking is required). Additionally, the number of rows (columns) is divisible by the number of available processes.

**Note 2:** you can rely on the implementation of the `la::dense_matrix` class, whose declaration is provided below. This class can store integer values and is essentially the same class we used during exercise sessions for managing matrices of doubles, with which you should be familiar.

**dense\_matrix.hpp:**

```
#ifndef DENSE_MATRIX_HH
#define DENSE_MATRIX_HH

#include <istream>
#include <vector>

namespace la // Linear Algebra
{
    class dense_matrix final
    {
    public:
        typedef std::vector<int> container_type;

        typedef container_type::value_type value_type;
        typedef container_type::size_type size_type;
        typedef container_type::pointer pointer;
        typedef container_type::const_pointer const_pointer;
        typedef container_type::reference reference;
        typedef container_type::const_reference const_reference;

    private:
        size_type m_rows = 0, m_columns = 0;
        container_type m_data;

        size_type
        sub2ind (size_type i, size_type j) const;

    public:
        dense_matrix (void) = default;

        dense_matrix (size_type rows, size_type columns,
                     const_reference value = 0.0);

        explicit dense_matrix (std::istream &);

        void
        read (std::istream &);
    };
}
```

```

void
swap (dense__matrix &);

reference
operator () (size__type i, size__type j);
const_reference
operator () (size__type i, size__type j) const;

size__type
rows (void) const;
size__type
columns (void) const;

dense__matrix
transposed (void) const;

pointer
data (void);
const_pointer
data (void) const;

void
print (std::ostream& os) const;
};

dense__matrix
operator * (dense__matrix const &, dense__matrix const &);

void
swap (dense__matrix &, dense__matrix &);
}

#endif // DENSE_MATRIX_HH

```

## Solution 2

1. To compute the desired result in a parallel fashion, we can proceed in the following way:
  - (i) We first determine if a single program argument is provided at the command line and accordingly generate a message.
  - (ii) If the input matrix file name is provided correctly, rank 0 reads the file and shards the matrix by row through a *scatter* operation.
  - (iii) Each process verifies that its matrix portion satisfies the diagonal condition (all non-diagonal elements are zero).
  - (iv) The global result is obtained by employing a *reduce* operation that computes the logical AND across all processes.
  - (v) Finally, rank 0 outputs the result.
2. A possible piece of code implementing the mechanisms above is the following:

```

1  #include <iostream>
2  #include <string>
3  #include <fstream>
4  #include <mpi.h>
5  #include "dense_matrix.hpp"
6
7  bool isDiagonalMatrix(const la::dense__matrix& matrix);
8

```

```

9  int main(int argc, char *argv[]) {
10     int rank, size;
11     MPI_Init(&argc, &argv);
12     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
13     MPI_Comm_size(MPI_COMM_WORLD, &size);
14
15     unsigned n; // Size of the matrix
16     la::dense_matrix matrix;
17
18     // (i)
19     // Check arguments number is correct and print instructions for
20     // program use
21     if (argc != 2) {
22         if (rank == 0)
23             std::cerr << "Usage:" << argv[0] << "<filename>"<<std::endl;
24         MPI_Finalize();
25         return -1;
26     }
27
28     // (ii-a)
29     // Argument number is ok. Let's read the matrix from file and share
30     // across all ranks
31     if (rank == 0) {
32         std::ifstream ifs(argv[1]);
33         matrix.read(ifs);
34         ifs.close();
35         // Print matrix
36         std::cout << "Input matrix:" << std::endl;
37         matrix.print(std::cout);
38     }
39
40     // (ii-b)
41     n = matrix.rows();
42     // Broadcast the size of the matrix to all processes
43     MPI_Bcast(&n, 1, MPI_UNSIGNED, 0, MPI_COMM_WORLD);
44
45     // Calculate the number of rows each process will handle
46     unsigned local_n = n / size;
47     la::dense_matrix sub_matrix(local_n,n,0);
48
49
50     // Scatter the matrix to all processes
51     MPI_Scatter(matrix.data(), local_n * n, MPI_INT, sub_matrix.data(),
52                local_n * n, MPI_INT, 0, MPI_COMM_WORLD);
53
54     // (iii)
55     // Each process checks its part of the matrix
56     unsigned local_result = isDiagonalMatrix(sub_matrix);
57     std::cout << "local_result, rank: " << rank << " " << local_result << std::endl;
58
59     // (iv)
60     if (rank == 0)
61         MPI_Reduce(MPI_IN_PLACE, &local_result, 1, MPI_UNSIGNED, MPI_LAND, 0,
62                   MPI_COMM_WORLD);
63     else
64         MPI_Reduce(&local_result, nullptr, 1, MPI_UNSIGNED, MPI_LAND, 0,
65                   MPI_COMM_WORLD);
66

```



```

65 // (v)
66 // Rank 0 print the result
67 if (rank == 0) {
68     if (local_result)
69         std::cout << "The input matrix is a diagonal matrix" << std::endl;
70     else
71         std::cout << "The input matrix is NOT a diagonal matrix" << std::endl;
72 }
73
74
75 MPI_Finalize();
76 return 0;
77 }
78
79
80
81
82
83 bool isDiagonalMatrix(const la::dense_matrix& matrix) {
84     int rank;
85     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
86     // Check for diagonal properties considering an index shift
87     for (int i = 0; i < matrix.rows(); ++i) {
88         for (int j = 0; j < matrix.columns(); ++j) {
89             if (i + rank * matrix.rows() != j && matrix(i,j) != 0) {
90                 return false; // Found a non-zero element outside the diagonal
91             }
92         }
93     }
94
95     return true; // All non-diagonal elements are zero
96 }

```

Through the `MPI_Init` function, the MPI environment is initialized. Subsequently, after performing argument verification, rank 0 reads the input file and broadcasts the matrix size ( $n$ ) to all processes using the `MPI_Bcast` function. Rows are distributed across all processes using the `MPI_Scatter` function. In this manner, each process checks, through the `isDiagonalMatrix` function, whether its matrix portion satisfies the diagonal condition. This function determines if a given matrix is diagonal by iterating over each element and checking if non-diagonal elements are zero. It employs a rank-based index shift ( $i + rank * local\_n$ ) to access elements.

The final result is computed by an `MPI_Reduce` logical AND (`MPI_LAND`) operation. The `MPI_IN_PLACE` option enables the root process (rank 0) to utilize the same buffer for both sending and receiving during reduction. Alternatively, despite requiring additional communication, all ranks can execute `MPI_Allreduce(MPI_IN_PLACE, &local_result, 1, MPI_UNSIGNED, MPI_LAND, MPI_COMM_WORLD)` to obtain the result, which is equivalent to lines 59—62.

### Exercise 3 (4 points)

Consider the following piece of code:

```
1  #include <iostream>
2
3  class BinaryIntOp {
4  public:
5      virtual int calc(int v1, int v2) = 0;
6  };
7
8  class Add : public BinaryIntOp {
9  public:
10     int calc(int v1, int v2) override {
11         return v1 + v2;
12     }
13 };
14
15 class Sub : public BinaryIntOp {
16 public:
17     int calc(int v1, int v2) override {
18         return v1 - v2;
19     }
20 };
21
22 class Mult : public BinaryIntOp {
23 public:
24     int calc(int v1, int v2) override {
25         return v1 * v2;
26     }
27 };
28
29 class Div : public BinaryIntOp {
30 public:
31     int calc(int v1, int v2) override {
32         return v1 / v2;
33     }
34 };
35
36 class OppAdd : public Add {
37 public:
38     int calc(int v1, int v2) override {
39         return -(Add::calc(v1, v2));
40     }
41 };
42
43 class AbsMult : public Mult {
44 public:
45     int calc(int v1, int v2) override {
46         return abs(Mult::calc(v1, v2));
47     }
48 };
49
```

```

50 int main() {
51
52     BinaryIntOp *op1;
53     Add op2;
54     Mult op3;
55     Div op4;
56     OppAdd op5;
57     AbsMult op6;
58
59     op1 = &op2;
60     int res1 = op1->calc(op2.calc(op3.calc(-2, 6), op4.calc(3, 7)), 1);
61
62     std::cout << res1 << "\n";
63
64     Mult &op7 = op6;
65     op1 = new Sub();
66     int res2 = op1->calc(op5.calc(-4, op7.calc(5, op4.calc(-4, 2))), op6.calc(-10, 2));
67
68     std::cout << res2 << "\n";
69
70     op2 = op5;
71     op1 = &op5;
72     int res3 = op2.calc(op1->calc(3, 7), op7.calc(-1, op2.calc(0, (new Sub())->calc(0, 8))));
73
74     std::cout << res3 << "\n";
75
76     op1 = &op7;
77     Mult &op8 = op3;
78     int res4 = op1->calc(2, (new Sub())->calc(op3.calc(10, op8.calc(5, -2)), 10));
79
80     std::cout << res4 << "\n";
81
82     return 0;
83 }

```

What numbers are printed at the following lines?

- (a) Line 62;
- (b) Line 68;
- (c) Line 74;
- (d) Line 80.

### Solution 3

TBC

- (a) **-11**; indeed, at line 60, where the result is computed, `op1` is of type `Add`, `op2` is of type `Add`, `op3` is of type `Mult`, and `op4` is of type `Div`; hence, the computation amounts to:  $((-2 * 6) + (3/7)) + 1 = -11$  (recall that `Div` is the division among integers, so  $3/7 = 0$ ).
- (b) **-26**; at line 66 `op1` is of type `Sub`, `op5` is of type `OppAdd`, `op7` is of type `AbsMult`, `op4` is of type `Div` `op6` is of type `AbsMult`; hence, the computation amounts to:  $-(-4 + |5 * (-4/2)|) - |-10 * 2| = -26$ .
- (c) **-2**; at line 72 `op2` is of type `Add`, `op1` is of type `OppAdd`, and `op7` is of type `AbsMult`; hence, the computation amounts to:  $-(3 + 7) + |-1 * (0 + (0 - 8))| = -2$ .
- (d) **220**; at line 78 `op1` is of type `AbsMult`, `op3` is of type `Mult`, and `op8` is of type `Mult`; hence, the computation amounts to:  $|2 * ((10 * (5 * -2)) - 10)| = 220$ .