

Streams & I/O

Federica Filippini, Luca Pozzoni,
Giovanni Quattrocchi, Danilo Ardagna

Politecnico di Milano

{nome.cognome}@polimi.it

Content

- Fundamental I/O concepts
- Files
 - Opening
 - Reading and writing streams
- File modes and Binary I/O
- String streams
 - Line-oriented input
- Examples

Files

- We turn our computers on and off
 - The contents of its main memory is transient
- Data needs to be preserved
 - It must be stored on disks and similar permanent storage devices
- A file is a sequence of bytes stored in permanent storage
 - A file has a name
 - The data on a file has a format
- We can read/write a file if we know its name and format

A file


0: 1: 2:



- At the fundamental level, a file is a **sequence of bytes** numbered from 0 upwards
- Other notions can be supplied by programs that interpret a “**file format**”
 - For example, the 6 bytes corresponding to "123.45" might be interpreted as the floating-point number 123.45

Files

- General model



Files

- **To read a file**
 - We must know its name
 - We must open it (for reading)
 - Then we can read it
 - Once finished, we must close it
 - That is typically done implicitly (when the stream object is destroyed)
- **To write a file**
 - We must name it
 - We must open it (for writing)
 - Or create a new file of that name
 - Then we can write it
 - Once finished, we must close it
 - That is typically done implicitly (when the stream object is destroyed)

Opening a file for reading

```
// ...  
int main()  
{  
    cout << "Please enter input file name: ";  
    string iname;  
    cin >> iname;  
  
    ifstream ist {iname}; // ifstream is an "input stream from a file"  
                        // defining an ifstream with a name string  
                        // opens the file of that name for reading  
  
    if (!ist) error("can't open input file ", iname);  
    // ...
```

Opening a file for writing

```
// ...  
cout << "Please enter name of output file: ";  
string oname;  
cin >> oname;
```

```
ofstream ofs {oname} ;// ofstream is an “output stream from a file”  
                        // defining an ofstream with a name string  
                        // opens the file with that name for writing
```

```
if (!ofs) error("can't open output file ", oname);  
// ...  
}
```


Implicit close

- When an **fstream** object goes out of scope, the file it is bound to is automatically closed

```
if (read) {  
    // create input and open the file  
    ifstream input{name};  
    if (input) { // if the file is ok, "process" this file  
        process(input);  
    } else  
        cerr << "couldn't open: " + ;  
    // input goes out of scope and is destroyed on each iteration  
}
```

Reading from a file

- Suppose a file contains a sequence of pairs representing hours and temperature readings

```
0    60.7  
1    60.6  
2    60.3  
3    59.22
```

- The hours are numbered from **0** to **23**
- No further format is assumed
 - Maybe we can do better than that (but not just now)
- Termination
 - Reaching the **end-of-file** terminates the read
 - Anything unexpected in the file terminates the read
 - E.g., character 'q'

Reading a file

```

struct Reading { // a temperature reading
    int hour; // hour after midnight [0:23]
    double temperature;
};

vector<Reading> temps; // create a vector to store the readings

int hour;
double temperature;
ifstream ist{fname};
while (ist >> hour >> temperature) { // read
    if (hour < 0 || 23 < hour)
        cout << "hour out of range" << endl;
    temps.push_back( Reading{hour, temperature} ); // store
}

```

```

ifstream ist{fname};
while (ist >> hour >> temperature) { . // read
    temps.push_back( Reading{hour, temperature} ); // store
}
    
```

```

| 0  60.7
q 60.6
2  60.3
3  59.22
    
```

```

0 | 60.7
q 60.6
2  60.3
3  59.22
    
```

```

0  60.7
| q 60.6
2  60.3
3  59.22
    
```

```

0  60.7
| q 60.6
2  60.3
3  59.22
    
```

No Copy or Assign for I/O Objects

- **We cannot copy or assign objects of the IO types:**

```
ofstream out1, out2;  
out1 = out2; // error: cannot assign stream objects ofstream  
ofstream print(out1); // error: can't initialize the ofstream param  
out2 = print(out1); // error: cannot copy stream objects
```

Because we can't copy the IO types, we cannot have a parameter or return type that is one of the stream types

- Functions that do IO typically pass and return the stream through **references**
- Reading or writing an IO object changes its state, so the reference **must not be const**

```
void print(outstream&); // OK
```

Use

```
void do_some_printing(Date d1, Date d2)
{
    cout << d1; // means:
                // operator<<(cout,d1) ;

    cout << d1 << d2;
        // means:
        // (cout << d1) << d2; same as:
        // (operator<<(cout,d1)) << d2; same as:
        // operator<<((operator<<(cout,d1)), d2);
}
```


File Modes and Binary I/O

File open modes

- By default, an ifstream opens its file for reading
- By default, an ofstream opens its file for writing.
- Alternatives:
 - `ios_base::app` // append (i.e., output adds to the end of the file)
 - `ios_base::ate` // “at end” (open and seek to end)
 - `ios_base::binary` // binary mode – beware of system specific behavior
 - `ios_base::in` // for reading
 - `ios_base::out` // for writing
 - `ios_base::trunc` // truncate file to 0-length
- A file mode is optionally specified after the name of the file:
 - `ofstream of1 {name1};` // defaults to `ios_base::out`
 - `ifstream if1 {name2};` // defaults to `ios_base::in`
 - `ofstream ofs {name, ios_base::app};` // append rather than overwrite
 - `fstream fs {"myfile", ios_base::in | ios_base::out};` // both in and out

Text vs. binary files

123 as
characters:



12345 as
characters:



123 as binary:



In binary files, we use
sizes to delimit values

12345 as
binary:



123456 as
characters:



In text files, we use
separation/termination
characters

123 456 as
characters:



Text vs. binary

- Use text when you can
 - You can read it (without a fancy program)
 - You can debug your programs more easily
 - Text is portable across different systems
 - Most information can be represented reasonably as text
- Use binary when you must
 - E.g. image files, sound files

String Streams

String streams

- A **stringstream** reads/writes from/to a **string** rather than a file or a keyboard/screen

```
double str_to_double(string s)
// if possible, convert characters in s to floating-point value
{
    stringstream is {s}; // make a stream so that
                          // we can read from s

    double d;
    is >> d;
    if (!is) error("double format error: ", s);
    return d;
}

double d1 = str_to_double("12.4");           // testing
double d2 = str_to_double("1.34e-3");
double d3 = str_to_double("twelve point three"); // error
```

String streams

- See textbook for **ostringstream**
- String streams are very useful for
 - formatting into a fixed-sized space (think GUI)
 - for extracting typed objects out of a string

Type vs. line

- Read a string

```
string name;
cin >> name;
cout << name << '\n';
```

// input: **Dennis Ritchie**

// output: **Dennis**

- Read a line

```
string name;
getline(cin, name);
cout << name << '\n';
// now what?
// maybe:
istringstream ss{name};
ss >> first_name;
ss >> second_name;
```

// input: **Dennis Ritchie**

// output: **Dennis Ritchie**

Examples

Example 1: reading a CSV file

- An `istreamstream` is often used when we have some work to do on an entire line, and other work to do with individual words within a line

Morgan,2015552368,8625550123

Drew,9735550130

Lee,6095550132,2015550175,8005550000

// members are public by default

```
struct PersonInfo {  
    string name;  
    vector<string> phones;  
};
```



```
vector<PersonInfo> people; // will hold all the records from the input
string line;
ifstream data("data.csv");
```

```
// read the input a line at a time until cin hits end-of-file (or another error)
while ( getline(data, line) ) {
    PersonInfo info; // create an object to hold this record's data
    istringstream record(line); // bind record to the line we just read

    // read the name
    // note that we are changing the delimiter of getline to “,”
    getline(record, info.name, ',') ;
    string phone;
    // read the phone numbers
    while (getline(record, phone, ',')){
        info.phones.push_back(phone); // and store them
    }
    people.push_back(info); // append this record to people
}
```

Example 2: A Word Transformation Map

Write a program that given one string, transforms it into another. The input to our program is two files. The first file contains **rules** that we will use **to transform** the **text** in the second file. Each rule consists of a word that might be in the input file and a phrase to use in its place.

word-transformation file:

```
y why  
r are  
u you
```

second file:

```
where r u
```

output file:

```
where are you
```

DEMO

An example: A Word Transformation Map

```
void word_transform(istream &map_file, istream &input)
{
    auto trans_map = buildMap(map_file);
    string text;
    while (getline(input, text)) {
        istringstream stream(text);
        string word;
        bool firstword = true;
        while (stream >> word) {
            if (firstword)
                firstword = false;
            else
                cout << " ";
            cout << transform(word, trans_map);
        }
        cout << endl;
    }
}
```

An example: A Word Transformation Map

```
map<string, string> buildMap(ifstream &map_file)
{
    map<string, string> trans_map;
    string key;
    string value;
    while (map_file >> key && getline(map_file,
                                      value))
        if (value.size() > 1)
            trans_map[key] = value.substr(1);
        else
            cout << "no rule for " + key << endl;
    return trans_map;
}
```

An example: A Word Transformation Map

```
const string &
transform(const string &s, const map<string,
                                     string> &m)
{
    auto map_it = m.find(s);
    if (map_it != m.cend())
        // if this word is in the transformation map
        return map_it->second;
    else
        return s;
}
```

Readings

User-defined output: operator<<()

- Usually **trivial**:

```
ostream& operator<<( ostream& os, const Date& d )
{
    return os << '(' << d.year()
           << ',' << d.month()
           << ',' << d.day() << ')';
}
```

- We often use several different ways of outputting a value
 - Tastes for output layout and detail vary

User-defined input: operator>>()

```
istream& operator>>(istream& is, Date& dd)  
    // Read date in format: year month day  
{  
    int y, d, m;  
    if (is >> y >> m >> d) {  
        dd = Date{y,m,d}; // update dd  
    }  
    return is;  
}
```


Binary files

```
int main()    // use binary input and output
{
    cout << "Please enter input file name\n";
    string iname;
    cin >> iname;
    ifstream ifs {iname, ios_base::binary}; // note: binary
    if (!ifs) error("can't open input file ", iname);

    cout << "Please enter output file name\n";
    string oname;
    cin >> oname;
    ofstream ofs {oname, ios_base::binary}; // note: binary
    if (!ofs) error("can't open output file ", oname);

    // “binary” tells the stream not to try anything clever operation
    // with the bytes
```

Binary files

```
vector<int> v;
```

```
// read from binary file:
```

```
for (int i; ifs.read(as_bytes(i), sizeof(int)); )
```

```
// note: reading bytes
```

```
    v.push_back(i);
```

```
// ... do something with v ...
```

```
// write to binary file:
```

```
for(int i=0; i<v.size(); ++i)
```

```
    ofs.write(as_bytes(v[i]), sizeof(int)); // note: writing  
                                                // bytes
```

```
return 0;
```

```
}
```


```
// For now, treat as_bytes() as a primitive
```

```
// Warning! Beware transferring between different systems
```

Positioning in a filestream

- To support random access, the system maintains a marker that determines where the next read or write will happen
- We also have two functions:
 - One **repositions the marker** by seeking to a given position
 - The second **tells the current position** of the marker
- The library actually defines two pairs of seek and tell functions:
 - One pair is used by input streams, the other by output streams
 - The input and output versions are distinguished by a suffix that is either a **g** (“getting”, i.e. reading data), or **p** (“putting”, i.e. writing data)

Positioning in a filestream



```
fstream fs {name}; // open for input and output
// ...
fs.seekg(5); // move reading position ( 'g' for 'get' ) to 5 (the 6th character)
char ch;
fs>>ch; // read the x and increment the reading position to 6
cout << "sixth character is " << ch << ' (' << int(ch)
    << " ) \n";
fs.seekp(1); // move writing position ( 'p' for 'put' ) to 1 (the 2nd character)
fs<<'y'; // write and increment writing position to 2
```

Positioning in a filestream

- We can use only the **g** versions on an **istream** and on the types that inherit from it, **ifstream** and **istringstream**
- We can use only the **p** versions on an **ostream** and on the types that inherit from it, **ofstream** and **ostringstream**
- An **iostream**, **fstream**, or **stringstream** can both **read** and **write** the associated stream; we can use either the **g** or **p** versions on objects of these types

There Is Only One Marker

- The fact that the library distinguishes between the “**putting**” and “**getting**” versions of the **seek** and **tell** functions can be misleading
- Even though the library makes this distinction, **it maintains only a single marker in a stream** — there are no distinct **read** and **write** markers

Repositioning the Marker

// set the marker to a fixed position

seekg(new_position) ; *// set the read marker to the given pos_type*
// location

seekp(new_position) ; *// set the write marker to the given pos_type*
// location

// offset some distance ahead of or behind the given starting point

seekg(offset, from) ; *// set the read marker offset distance from*
from **seekp(offset, from) ;** *// offset has type off_type*


Reading and writing to the same file

abcd

efg

hi

j



abcd

efg

hi

j

5 9 12 14

Reading and writing to the same file

```
int main()
{
    // open for input and output and reposition file pointers to end-of-file
    // file mode argument
    fstream inOut("copyOut", fstream::ate | fstream::in |
                  fstream::out);

    if (!inOut) {
        cerr << "Unable to open file!" << endl;
        return EXIT_FAILURE; //EXIT_FAILURE
    }

    // inOut is opened in ate mode, so it starts out positioned at the end
    auto end_mark = inOut.tellg(); // remember original end-of-file
                                   // position
    inOut.seekg(0, fstream::beg); // reposition to the start of the file
    size_t cnt = 0; // accumulator for the byte count string line;
    string line; // hold each line of input
```

Reading and writing to the same file

```

// while we haven't hit an error and are still reading the original
// data and can get another line of input
while (inOut && inOut.tellg() != end_mark
      && getline(inOut, line))
{
    cnt += line.size() + 1; // add 1 to account for the newline
    auto mark = inOut.tellg(); // remember the read position
    inOut.seekp(0, fstream::end); // set the write marker to the end
    inOut << cnt; // write the accumulated length
    // print a separator if this is not the last line
    if (mark != end_mark) inOut << " ";
    inOut.seekg(mark); // restore the read position
}
inOut.seekp(0, fstream::end); // seek to the end
inOut << "\n"; // write a newline at end-of- file
return 0;
}

```

Positioning

- Whenever you can
 - Use simple streaming
 - Streams/streaming is a very powerful metaphor
 - Write most of your code in terms of “plain” istream and ostream
 - Positioning is far more error-prone
 - Handling of the end of file position is system dependent and basically unchecked

Using ostringstreams

- An **ostringstream** is useful when we need to build up our output a **little at a time** but **do not** want to print the output **until later**
- For example, we might want to validate and reformat the phone numbers we read in the previous example
 - If all the numbers are valid, we want to print a new file containing the reformatted numbers
 - If a person has any invalid numbers, we won't put them in the new file. Instead, we'll write an error message containing the person's name and a list of their invalid numbers

Using ostringstreams

```

for (const auto &entry : people) { // for each entry in people
    ostringstream formatted, badNums; // objects created on each loop
    for (const auto &nums : entry.phones) { // for each number
        if (!valid(nums)) {
            badNums << " " << nums; // string in badNums
        } else
            // "writes" to formatted's string
            formatted << " " << format(nums);
        }
    if (badNums.str().empty()) // there were no bad numbers
        os << entry.name << " " // print the name
        << formatted.str() << endl; // and reformatted numbers
    else
        // otherwise, print the name and bad numbers
        cerr << "input error: " << entry.name
        << " invalid number(s) " << badNums.str() << endl;
}

```

References

- Lippman Chapters 8, 17

Credits

- Bjarne Stroustrup. www.stroustrup.com/Programming