

Standard Template Library 2

Mehrnoosh Askarpour & Danilo Ardagna

Politecnico di Milano

name.lastname@polimi.it

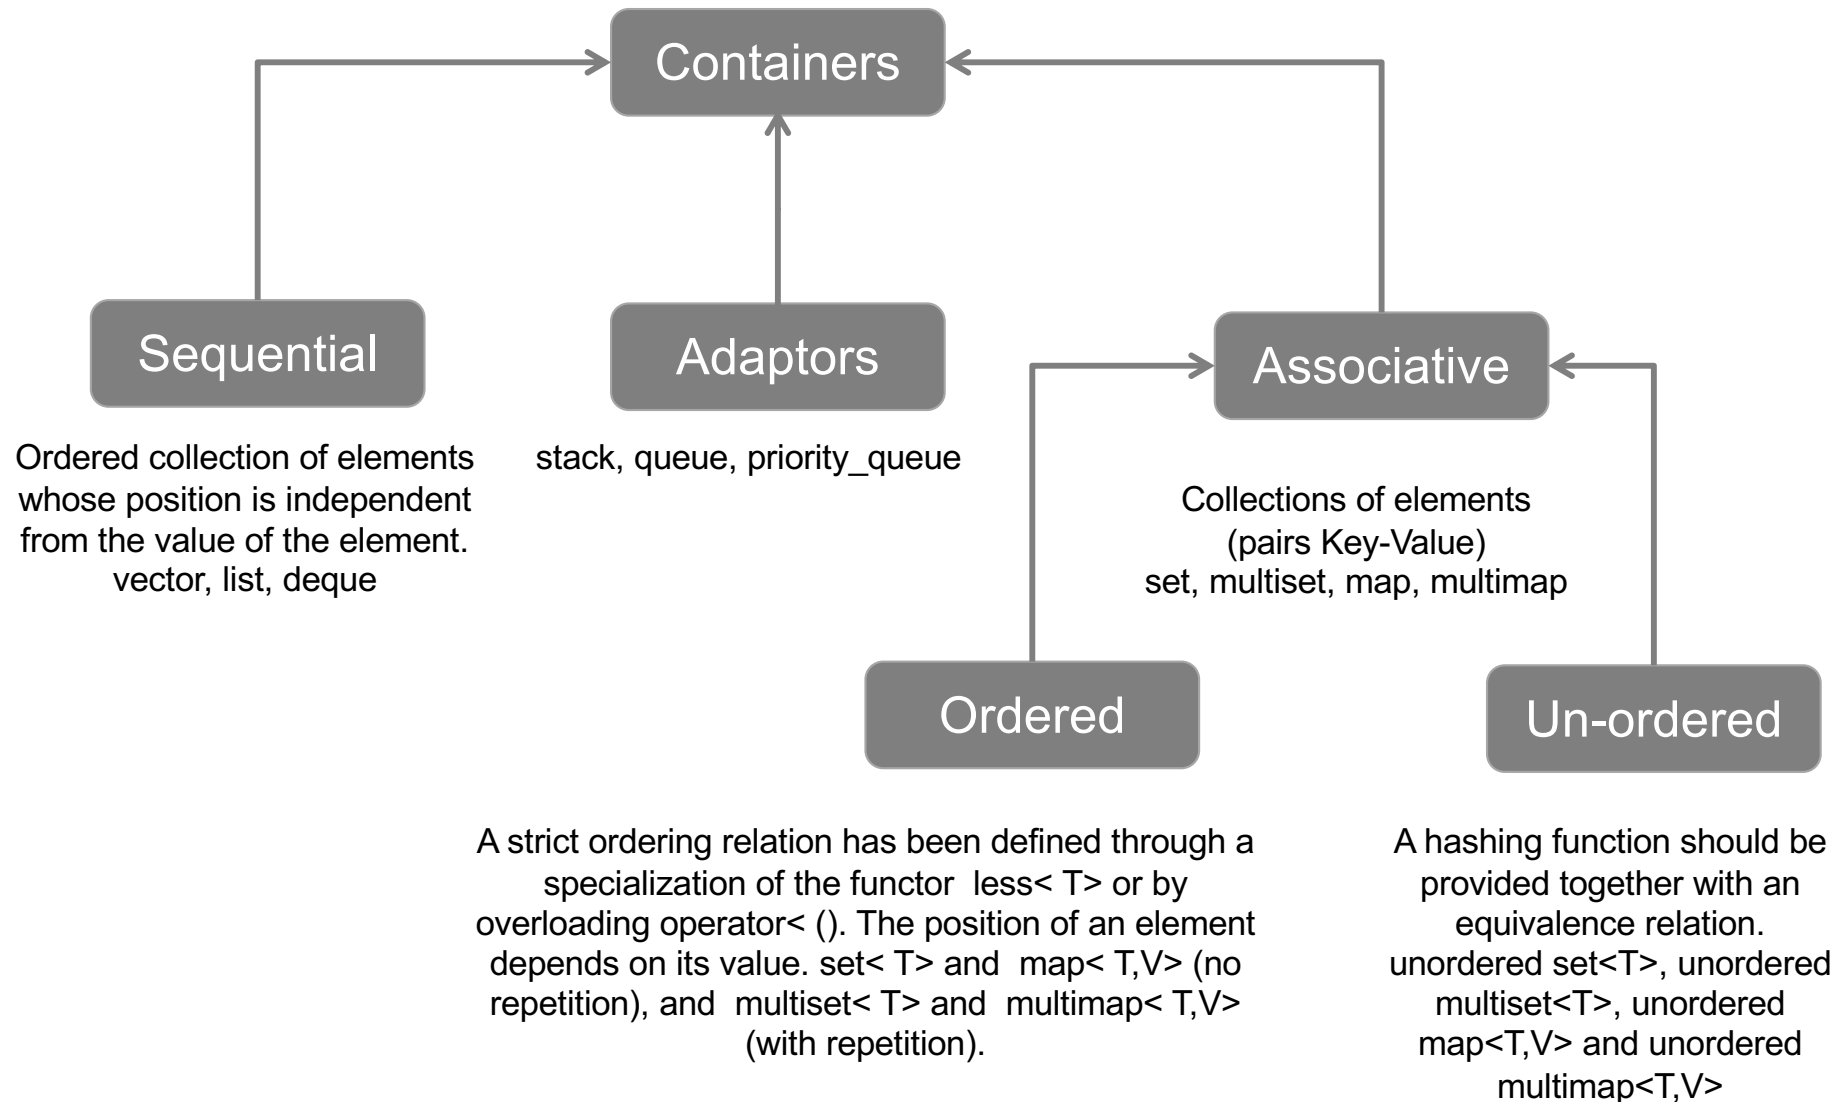


**POLITECNICO
DI MILANO**

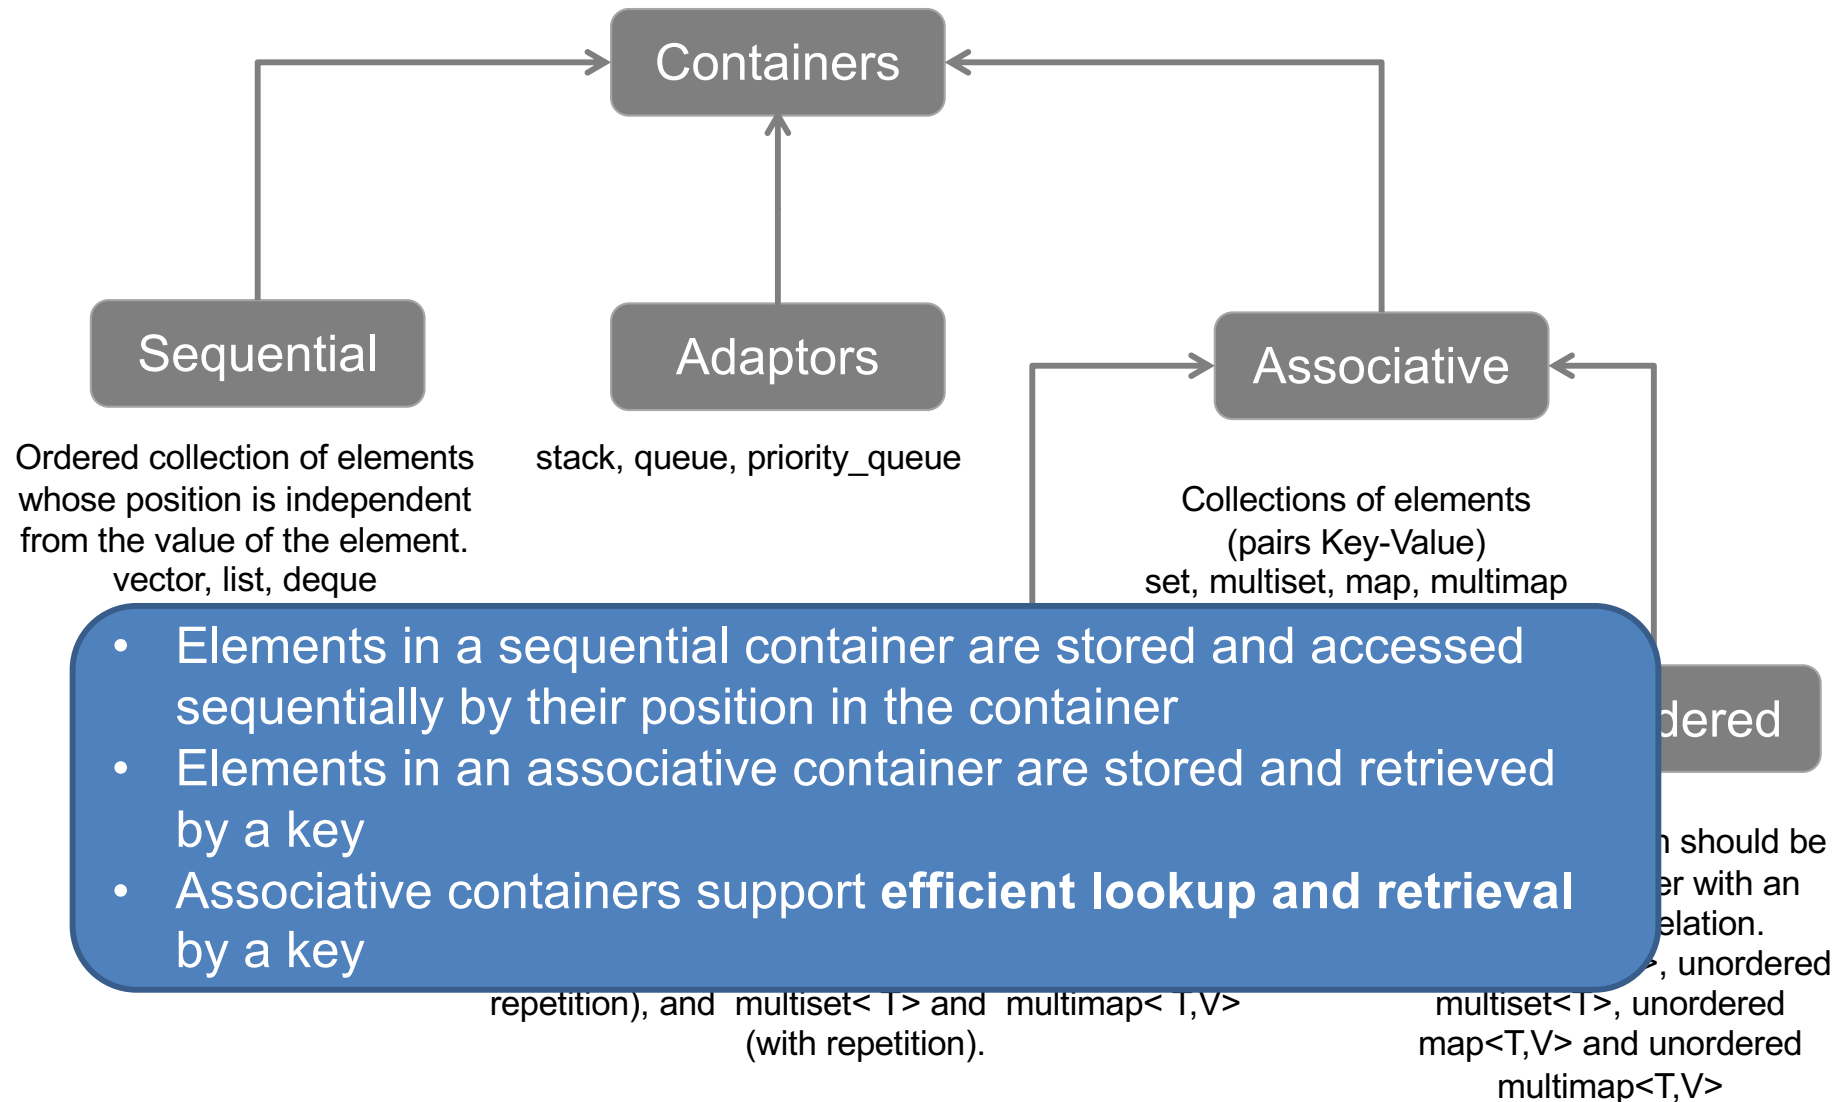
Content

- Adaptors
- Associative Containers
 - set, map
 - multiset, multimap
 - Unordered map, set

Sequential and Associative containers



Sequential and Associative containers



Adaptors

Container Adaptors (readings)

- Container adaptors are interfaces created on top of a (limited) set of functionalities of a pre-existing sequential container, which provide a different API
- When you declare the container adapters, you have an option of specifying which sequential container to use as underlying container

Container Adaptors (readings)

- `stack`:
 - Container providing Last-In, First-Out (LIFO) access
 - You remove (**pop**) elements in the reverse order you insert (**push**) them. You cannot get any elements in the middle
 - Usually this goes on top of a **deque**
- `queue`:
 - Container providing First-In, First-Out (FIFO) access
 - You remove (**pop**) elements in the same order you insert (**push**) them. You cannot get any elements in the middle but only **front** and **back**
 - Usually this goes on top of a **deque**
- `priority_queue`:
 - Container providing sorted-order access to elements
 - You can insert (**push**) elements in any order, and then retrieve (**pop**) the "highest priority" of these values at any time
 - Priority queues in C++ STL use a **heap** structure internally, which in turn is basically **array-backed**; thus, usually this goes on top of a **vector**

Associative Containers

Overview of of the Associative Containers

Associative containers support the **general container operations**. However, they **do not support** the **sequential-container position-specific** operations, such as `push_front`. Because the elements are stored based on their **keys**, these operations would be meaningless for the associative containers.

Nevertheless they have:

- Type aliases
- **Bidirectional iterators**
- Specific operations
- Hash functions (unordered version)

Overview of of the Associative Containers

- `map`: holds key-value pairs
 - `multimap`: a map in which one key can appear multiple times
- } `<map>`
- `set`: the key is the value
 - `multiset`: a set in which a key can appear multiple times
- } `<set>`
- `unordered-map`: a map organized by a hash function
 - `unordered-multimap`: hashed map, keys can appear multiple times
- } `<unordered_map>`
- `unordered-set`: a set organized by a hash function
 - `unordered-multiset`: hashed set, keys can appear multiple times
- } `<unordered_set>`

Consider multi- version as readings

std::map

- A `map` is a collection of `<key, value>` pairs with **unique keys**
- It is often referred to as an **associative array**
 - An associative array is like a “normal” array except that its **subscripts don't have to be integers**
- Values in a `map` are found by a key rather than by their position
- Example: Given a map of names to phone numbers, we'd use a person's name as a subscript to fetch that person's phone number:
 - E.g.: `cout << phone_numbers["Mario"];`

std::map

// count the number of times each word occurs

// in the input

```
map<string, size_t> word_count;
```

```
string word;
```

```
while (cin >> word) // fetch and increment the  
                        //counter for word
```

```
    ++word_count[word];
```

```
for (const auto &w : word_count)
```

```
    // for each element in the map print the results
```

```
        cout << w.first << " occurs " <<  
            w.second<<
```

```
            ((w.second > 1) ?
```

```
            " times" : " time") << endl;
```

std::map

// count the number of times each word occurs

// in the input

```
map<string, size_t> word_count;
```

```
string word;
```

```
while (cin >> word) // fetch and increment the  
                        //counter for word
```

```
    ++word_count[word];
```

```
for (const auto &[key, value] : word_count)
```

```
    // for each element in the map print the results
```

```
        cout << key << " occurs " <<  
              value <<
```

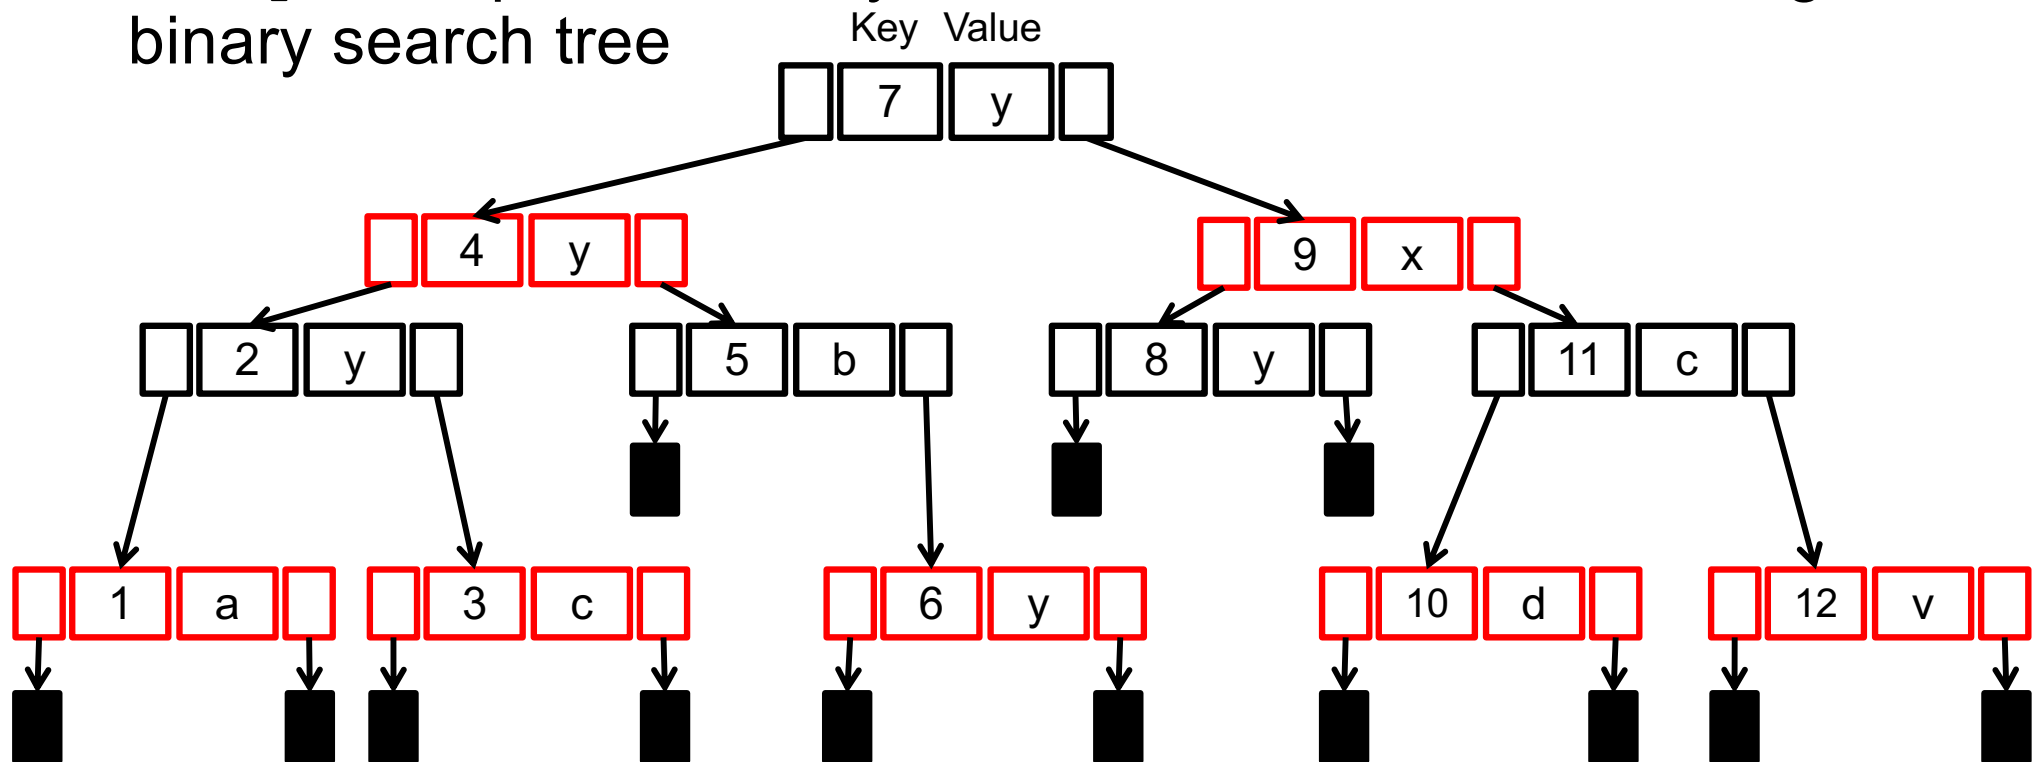
```
            ((value > 1) ?
```

```
            " times" : " time") << endl;
```

New functionality since C++17

Implementation of a `map`

A `map` is implemented by red-black trees, self-balancing binary search tree



Insert and delete $O(\log n)$ at worst case. The rules keep the tree balanced

std::set

- A `set` is simply a collection of objects. A `set` is most useful when we simply want to know whether a value is present
- Example: a business might define a set named `bad_checks` to hold the names of individuals who have written bad checks. Before accepting a check, that business would query `bad_checks` to see whether the customer's name was present

- E.g.:

```
if (bad_checks.find("Mario") == bad_checks.end())  
    cout << "Mario is ok!"  
else  
    cout << "Mario is bad guy!";
```

std::set

- Example: excluding certain words from the map of word occurrences

// count the number of times each word occurs

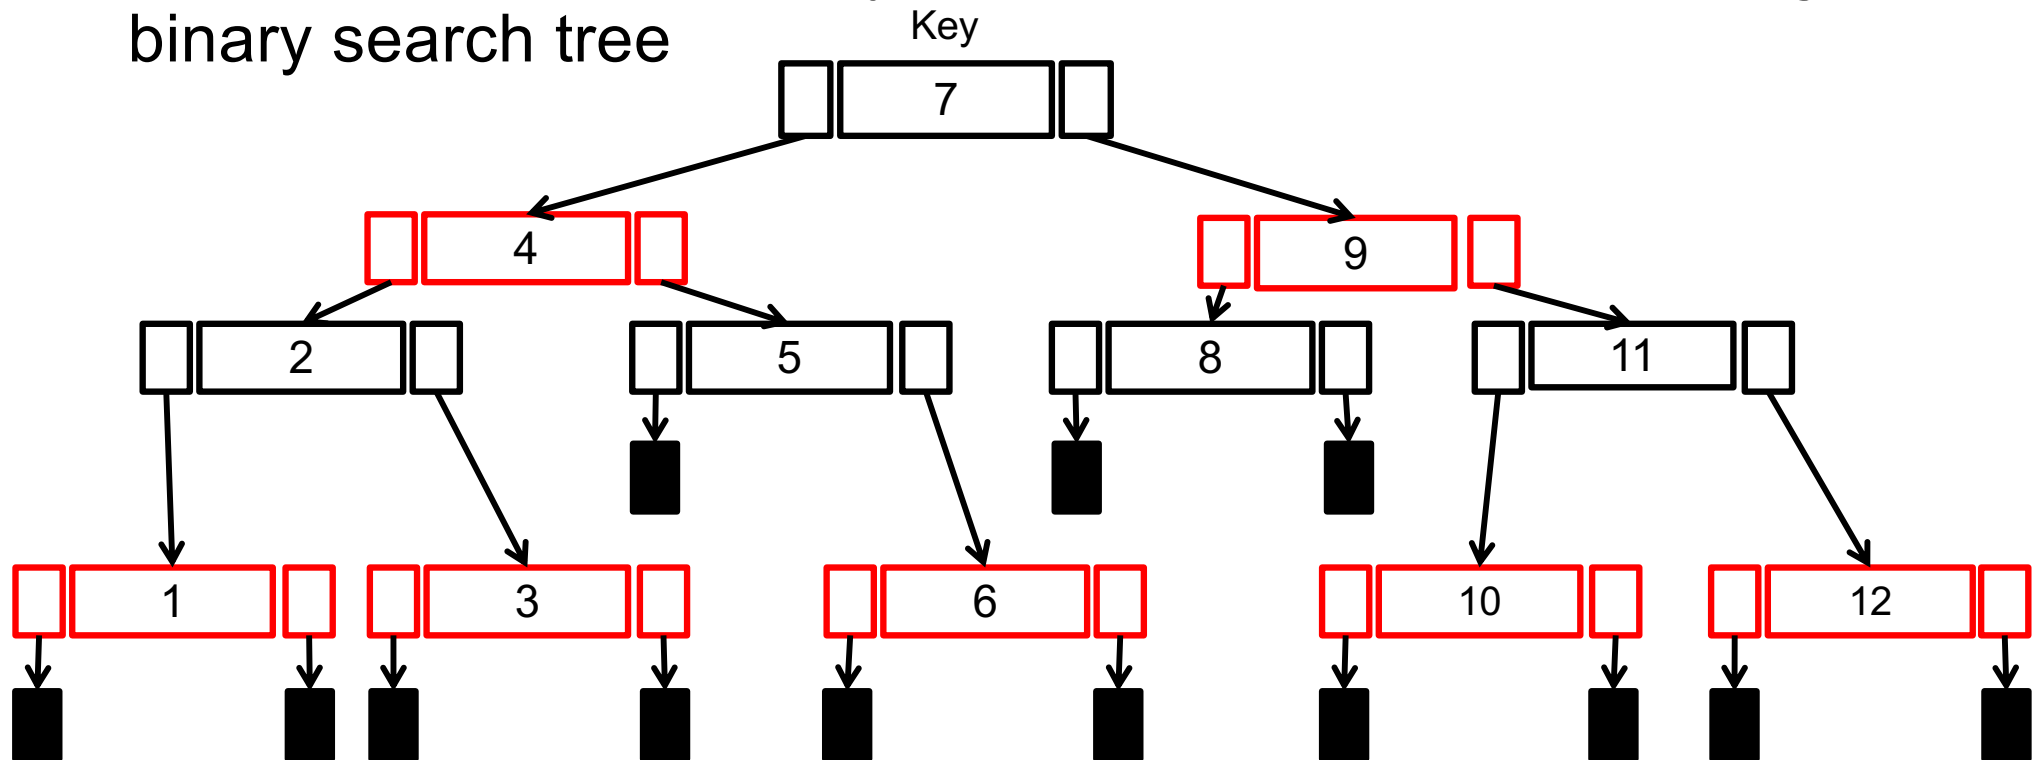
// in the input

```
map<string, size_t> word_count;
set<string> exclude = {"The", "But", "And", "Or",
                      "An", "A", "the", "but",
                      "and", "or", "an", "a"};

string word;
while (cin >> word)
    if (exclude.find(word) == exclude.end())
        ++word_count[word];
```


Implementation of a `set`

A `set` is implemented by red-black trees, self-balancing binary search tree



Insert and delete $O(\log n)$ at worst case. The rules keep the tree balanced

Set and Multiset - Example

```
vector<int> ivec;
for (vector<int>::size_type i = 0; i != 10; ++i) {
    ivec.push_back(i); // duplicate copies of each
    ivec.push_back(i); // number
}
set<int> iset(ivec.cbegin(), ivec.cend());
multiset<int> miset(ivec.cbegin(), ivec.cend());

cout << ivec.size() << endl; // prints 20
cout << iset.size() << endl; // prints 10
cout << miset.size() << endl; // prints 20
```

Associative vs. Sequential containers

- The associative containers do not support
 - **position-specific** operations, e.g., `push_front` or `back`
 - constructors or insert operations that **take an element value and a count**
- The associative container **iterators** are always **bidirectional**
 - Accessing `begin()` and `end()` is **$O(1)$**
 - `++it` and `--it` are **$O(1)$**

Requirements on Key Type

For the **ordered containers** the key type must define a way to compare the elements:

- By default, the library uses the \leq to compare the keys
- We can also supply our own $<$ operation to use a **strict weak ordering** over the key type
 - Two keys cannot both be “less than” each other
 - If k_1 is “less than” k_2 and k_2 is “less than” k_3 , then k_1 must be “less than” k_3
 - If there are two keys, and neither key is “less than” the other, then we’ll say that those keys are “equivalent.” If k_1 is “equivalent” to k_2 and k_2 is “equivalent” to k_3 , then k_1 must be “equivalent” to k_3

The `pair` type

A `pair` is a library type, defined in the `utility` header, which holds two data members.

// holds two strings

```
pair<string, string> anon;
```

// holds a string and an `size_t`

```
pair<string, size_t> word_count;
```

// holds string and `vector<int>`

```
pair<string, vector<int>> line;
```

```
pair<string, string> author{"James", "Joyce"};
```

The `pair` type

- The data members of `pair` are public
- These members are named `first` and `second`
- Elements in a `map` are `pairs`
- Only a limited number of operations defined in the library

<pre>pair<T1, T2> p; pair<T1, T2> p(v1, v2); pair<T1, T2> p={v1, v2};</pre>	Pair definition with or without initialization
<pre>make_pair(v1, v2)</pre>	Pair definition. Type of pair is inferred from v1 and v2 type.
<pre>p.first / p.second</pre>	Returns first or second member of p
<pre>p1(<, >, <=, >=, ==, !=) p2</pre>	Relational operators and equality. For example <code>p1<p2</code> if <code>p1.first<p2.first</code> or <code>!(p2.first<p1.first) && p1.second<p2.second</code>

Associative Container Type Aliases

- `key_type`: type of the key of the container
- `mapped_type`: type associated with each **map** key
- `value_type`: it is the same as `key_type` for sets and `pair<const key_type, mapped_type>` for maps.

Associative Container Type Aliases

- `key_type`: type of the key of the container
- `mapped_type`: type associated with each **map** key
- `value_type`: it is the same as `key_type` for sets and `pair<const key_type, mapped_type>` for maps.

```
set<string>::value_type v1;  
set<string>::key_type v2;  
map<string, int>::value_type v3;  
  
map<string, int>::key_type v4;  
map<string, int>::mapped_type v5;
```



Associative Container Type Aliases

- `key_type`: type of the key of the container
- `mapped_type`: type associated with each **map** key
- `value_type`: it is the same as `key_type` for sets and `pair<const key_type, mapped_type>` for maps.

```
set<string>::value_type v1; // v1 is a string
```

```
set<string>::key_type v2; // v2 is a string
```

```
map<string, int>::value_type v3; // v3 is a
```

```
                // pair<const string, int>
```

```
map<string, int>::key_type v4; // v4 is a string
```

```
map<string, int>::mapped_type v5; // v5 is an int
```

key is const!

Remember that the `value_type` of a `map` is a pair and that we can change the value but not the key member of that pair

```
auto map_it = word_count.begin();  
cout << map_it->first <<" " << map_it->second;
```

// error: key is const

```
map_it->first = "new key";
```

key is const!

The keys in a `set` are also `const`. We can use a `set` iterator to read, but not write, an element's value

```
set<int> iset = {0,1,2,3,4,5,6,7,8,9};  
set<int>::iterator set_it = iset.begin();  
if (set_it != iset.end()) {  
    // error: keys in a set are read-only  
    *set_it = 42;  
    // ok: can read the key  
    cout << *set_it << endl;  
}
```

Iterating across an associative container

When we use an iterator to traverse a map, multimap, set, or multiset, the iterators yield elements in **ascending key order**

```
auto map_it = word_count.cbegin();  
while (map_it != word_count.cend()) {  
    cout << map_it->first << " occurs "  
    << map_it->second << " times" << endl;  
    ++map_it;  
}
```

Adding Elements

Because (unordered) `map` and (unordered) `set` contain unique keys, inserting elements that are already present, has no effect

<code>c.insert(v)</code> <code>c.emplace(args)</code>	<code>v</code> <code>value_type</code> object. <i>args</i> are used to construct an element. Insert in map or set only if an element with the given key already is not in <code>c</code> . Return a pair of an iterator referring to the element with the given key and a bool indicating whether the element was inserted. For <code>multimap</code> and <code>multiset</code> it returns an iterator to the new element.
<code>c.insert(b, e)</code>	<code>b</code> and <code>e</code> are iterators denoting a range of <code>c::value_type</code> elements
<code>c.insert(il)</code>	<code>il</code> is a braced list of values
<code>c.insert(p, v)</code> <code>c.emplace(p, args)</code>	Like the first two, but uses <code>p</code> as a hint for where to begin the search for where the new element should be stored. Returns an iterator to the element with the given key

Adding Elements

Because (unordered) `map` and (unordered) `set` contain unique keys, inserting elements that are already present, has no effect

```
vector<int> ivec = {2,4,6,8,2,4,6,8}; // ivec has  
                                              // eight elements  
set<int> set2; // empty set  
set2.insert(ivec.cbegin(), ivec.cend()); // set2 has  
                                              // four elements  
set2.insert({1,3,5,7,1,3,5,7}); // set2 now has  
                                  // eight elements
```

Adding Elements

```
multimap<string, string> authors;
```

```
// adds the first element with the
```

```
// key Barth, John
```

```
authors.insert({"Barth, John", "Sot-Weed  
Factor"});
```

```
// ok: adds the second element with the key
```

```
// Barth, John
```

```
authors.insert({"Barth, John", "Lost in the  
Funhouse"});
```

Here `insert` returns always only an iterator to
the inserted element

Erasing Elements

We can erase one element or a range of elements by passing `erase` an iterator or an iterator pair.

<code>c.erase(k)</code>	Removes every element with key <code>k</code> from <code>c</code> . Returns <code>size_type</code> indicating the number of removed elements
<code>c.erase(p)</code>	Removes the element denoted by the iterator <code>p</code> . Returns an iterator to the element after <code>p</code>
<code>c.erase(b, e)</code>	Removes elements in the range from <code>b</code> to <code>e</code> , and returns <code>e</code>

Subscripting a map

- The `map` and `unordered_map` containers provide the **subscript operator** and a corresponding `at` function. The **set types do not support subscripting** because there is no “value” associated with a key in a set
- **We cannot subscript a `multimap` or an `unordered_multimap`** because there may be more than one value associated with a given key

<code>c[k]</code>	Returns the element with key <code>k</code> ; if <code>k</code> is not in <code>c</code>, adds a new, value initialized element with key <code>k</code>.
<code>c.at(k)</code>	Checked access to the element with key <code>k</code> ; <code>out_of_range</code> error if <code>k</code> is not in <code>c</code> .

Subscripting a map

```
map <string, size_t> word_count;  
// insert a value-initialized element with  
// key Anna; then assign 1 to its value  
word_count["Anna"] = 1;
```

Subscripting a map

```
map<string, size_t> word_count;  
// insert a value-initialized element with  
// key Anna; then assign 1 to its value  
word_count["Anna"] = 1;
```

- word_count is searched for the element whose key is Anna
- The element is not found
- A new key-value pair is inserted into word_count
 - The key is a const string holding Anna
 - The value is value initialized, i.e. it takes 0
- The newly inserted element is fetched set to 1

Subscripting a map

```
map <string, size_t> word_count;  
// insert a value-initialized element with  
// key Anna; then assign 1 to its value  
word_count["Anna"] = 1;
```

- Another way:

```
word_count.insert(make_pair("Anna",1));
```

Subscripting a map

```
map <string, size_t> word_count;  
// insert a value-initialized element with  
// key Anna; then assign 1 to its value  
word_count["Anna"] = 1;  
//fetch the element indexed by Anna  
cout << word_count["Anna"];  
// fetch the element and add 1 to it  
++word_count["Anna"];  
// fetch the element and print it  
  
cout << word_count["Anna"];
```

DEMO

Accessing Elements

<code>c.find(k)</code>	Returns an iterator to the first element with key k , or the off-the-end iterator if k is not in the container
<code>c.count(k)</code>	Returns the number of elements with key k. For the containers with unique keys, the result is always zero or one
<code>c.lower_bound(k)</code>	Return an iterator to the first element with key not less than k
<code>c.upper_bound(k)</code>	Return an iterator to the first element with key greater than k

Lower and upper bound are not valid for unordered containers, in that case you can rely on `equal_range`

lower and upper_bound

```
#include <iostream>
#include <map>
int main () {
    std::map<char,int> mymap;
    std::map<char,int>::iterator itlow,itup;
    mymap['a']=20;
    mymap['b']=40;
    mymap['c']=60;
    mymap['d']=80;
    mymap['e']=100;
    itlow=mymap.lower_bound ('b'); // itlow points to b
    itup=mymap.upper_bound ('d'); // itup points to e (not d!)
    mymap.erase(itlow,itup); // erases [itlow,itup)
    // print content:
    for (std::map<char,int>::iterator it=mymap.begin(); it!=mymap.end(); ++it)
        std::cout << it->first << " => " << it->second << '\n';
    return 0;
}
```

Accessing Elements

```
set<int> iset = {0,1,2,3,4,5,6,7,8,9};  
auto it = iset.find(1); // returns an iterator that  
                        // refers to the element with key 1  
it = iset.find(11); // returns the  
                  // iterator == iset.end()  
iset.count(1); // returns 1  
iset.count(11); // returns 0
```


GoodReads
<ul style="list-style-type: none"> - books [] - reviews []
<ul style="list-style-type: none"> + add_book(const string & title, unsigned pageNumber, const string &publisher, const string &author) + add_review(const string &bookTitle, const string &text, unsigned int rating) + get_avg_rating() + get_avg_rating(const string & title) + search_reviews(const vector<string> & keywords) + print_book(const string & title) - find_book(const string &title) - includes_all(const vector<string> &words, const vector<string> &keywords) - includes_word (const vector<string> &words, const string &k)

Book
<ul style="list-style-type: none"> - ratings_distr[] - pages_number - publisher - review_count - author - title; - avg_rating; - list<unsigned> review_indexes
<ul style="list-style-type: none"> + get_avg_rating() + add_review(unsigned index, unsigned stars) + to_string() + get_review_indexes() + get_title() - compute_rating()

Review
<ul style="list-style-type: none"> - book_title - text - rating - words []
<ul style="list-style-type: none"> + to_string() + get_text() + get_words() - find_in_words(const string & w)

- How to improve GoodReads::find_book() worst case complexity? $O(n_{\text{books}})$!
- How to simplify our code in GoodReads::search_review()?

GoodReads
<div> - map<string, BookData> books - reviews [] </div>
+ add_book(const string & title, unsigned pageNumber, const string &publisher, const string &author) + add_review(const string &bookTitle, const string &text, unsigned int rating) + get_avg_rating() + get_avg_rating(const string & title) + search_reviews(const vector<string> & keywords) + print_book(const string & title)

BookData
- ratings_distr[] - pages_number - publisher - review_count - author - avg_rating - list<unsigned> review_indexes
+ get_avg_rating() + add_review(unsigned index, unsigned stars) + to_string() + get_review_indexes() + get_title() - compute_rating()

Review
- book_title - text - rating - set<string> words
+ to_string() + get_text() + get_words() - find_in_words(const string & w)

- Introduce a map for books
- Drop the title in Book, i.e., refactor into BookData class
- Use sets for Review::words

A running example - GoodReads

```
class GoodReads {  
    map<string, BookData> books; // <title, BookData>  
    vector<Review> reviews;  
public:  
    void add_book(const string & title, unsigned pageNumber,  
                  const string &publisher,  
                  const string &author);  
    void add_review(const string &bookTitle,  
                    const string &text, unsigned int rating);  
    float get_avg_rating() const;  
    float get_avg_rating(const string & title) const;  
    void search_reviews(const vector<string> & keywords) const;  
    void print_book(const string & title) const;  
};
```

A running example - GoodReads

```
class BookData {
    vector<unsigned> ratings_distr;
    unsigned pages_number;
    string publisher;
    unsigned review_count;
    string author;
    float avg_rating;
    list<unsigned> review_indexes;
public:
    BookData(unsigned int pageNumber, const string
&publisher, const string &author);
    float get_avg_rating() const;
    void add_review(unsigned index, unsigned stars);
    string to_string() const;
    list<unsigned> get_review_indexes() const ;
private:
    float compute_rating();
};
```

A running example - GoodReads

```
class Review {
    string book_title;
    string text;
    unsigned rating;
    set<string> words;

public:
    Review(const string &bookTitle, const
string &text, unsigned int rating);
    string to_string() const;

    string get_text() const;
    set<string> get_words() const;

};
```

A running example - GoodReads

- `find_book` is not needed anymore, book search performed as:

```
const auto it = books.find(title);
```

- and if `it == books.cend()` the book is not in the collection otherwise it can be accessed through `it`

Worst case complexity now becomes $O(\log(n_books))$

A running example - GoodReads

- How to simplify keywords search in `search_reviews()`
- Given two sets `a` and `b` you can test if $a \supseteq b$ by:

```
std::includes(a.cbegin(), a.cend(),  
              b.cbegin(), b.cend())
```

- **where** `std::includes()` is defined in `algorithm` and returns a `bool`

A running example - GoodReads

```
void GoodReads::search_reviews(const vector<string> & keywords)
const {

    const set<string>keywords_set(keywords.cbegin(),keywords.cend());

    for (auto it = reviews.cbegin(); it != reviews.cend(); ++it){
        const set<string> & words = it->get_words();

        if(std::includes(words.cbegin(), words.cend(),
                        keywords_set.cbegin(), keywords_set.cend()))
            cout << it->to_string()<<endl;
    }

}
```


GoodReads method complexity

Vector based

- `Book::add_review()`
 - `push_back` in a vector
 - Worst case complexity $O(n_reviews)$
- `GoodReads::find_book()`
 - Sequential search in a vector
 - Worst case complexity $O(n_books)$
- `GoodReads::add_book()`
 - `find_book` and `push_back` in a vector
 - Worst case complexity $O(n_books)$
- `GoodReads::get_avg_rating()`
 - Sequential access in a vector
 - Worst case complexity $O(n_books)$
- `GoodReads::get_avg_rating(const string & title)`
 - `find_book`
 - Worst case complexity $O(n_books)$

Optimized version

- `Book::add_review()`
 - `push_back` in a list
 - Worst case complexity $O(1)$
- `GoodReads::find_book()`
 - `find` in a map
 - Worst case complexity $O(\log(n_books))$
- `GoodReads::add_book()`
 - `find` and `insert` in a map
 - Worst case complexity $O(\log(n_books))$
- `GoodReads::get_avg_rating()`
 - Sequential access in a map
 - Worst case complexity $O(n_books)$
- `GoodReads::get_avg_rating(const string & title)`
 - `find` in a map
 - Worst case complexity $O(\log(n_books))$

unordered_map &
unordered_set

Computer scientists' dream


- We love vectors
- We would like to have a bijective function such that given a key k in D we can compute the index i to store the value v associated to k :

$$f: D \rightarrow N$$

- That's just a dream



Computer scientists' dream



- We love vectors
- We would like a key k in D with v associated to it
- That's just a dream

$f: D \rightarrow V$

given
value

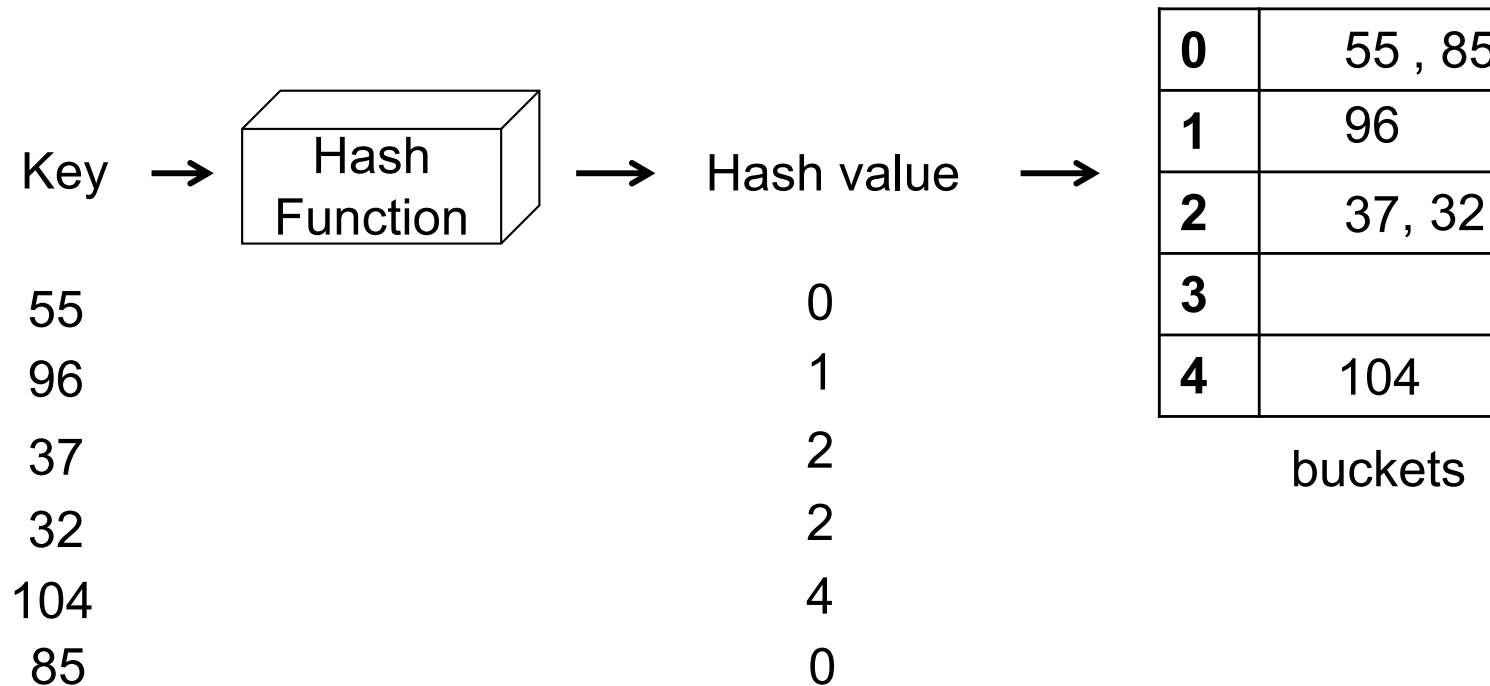
	surjective	non-surjective
injective	<p>bijjective</p>	<p>injective-only</p>
non-injective	<p>surjective-only</p>	<p>general</p>

From Wikipedia

Unordered Associative Containers



- Collection of buckets
 - Each bucket contains a variable number of items
- Use of a **hash function** to map elements to buckets
 - Given the item key, identify the proper bucket to store such item
 - All of the elements with a given hash value are stored in the same bucket
 - All the elements with the same key (in the multi- version) will be in the same bucket

Unordered Associative Containers



Different keys with the same hash value are stored in the same bucket and originate a collision

Unordered Associative Containers



Unordered Associative Containers

Key




Unordered Associative Containers

- The performance of an unordered container depends on the quality of its hash function and on the number and size of its buckets:
 - Average complexity $O(1)$
 - Worst case complexity $O(N)$
- **Rather than using a comparison operation** to organize their elements, these containers use a **hash function** and the **key type's `==` operator**
- Use an unordered container if the key type is inherently unordered or if performance testing reveals problems that hashing might solve

DEMO

A Simple Hash Function to Manage an `unordered_map`



```
unsigned my_hash_func(unsigned x, unsigned size)
{
    return x % size;
}
```

In APSC you will see how to define your hash functions. In APC we will rely on the ones defined for the **built-in** and **STL types**

Using an Unordered Container

```
unordered_map<string, size_t> word_count;
string word;
while (cin >> word)
    ++word_count[word];
for (const auto &w : word_count)
    cout << w.first << " occurs " <<
    w.second << ((w.second > 1) ? "
    times" : " time") << endl;
```

*// count occurrences, but the words won't be
// in alphabetical order*

Complexities of operations

You can choose the most suitable container in terms of complexity, depending on what operations you need to apply on them

Container	Operation		
	Insert	Find	Delete
list/ forward_list	$O(1)$	$O(n)$	$O(1)$
set/map	$O(\log(n))$	$O(\log(n))$	$O(\log(n))$
unordered set/map	$O(1)$ or $O(n)$	$O(1)$ or $O(n)$	$O(1)$ or $O(n)$

map **VS.** unordered_map

- If the most frequent operations of your application are **find**, **insert** or **delete** of single elements and you access through a key:
 - Use a `map` if you want to optimize the worst case complexity
 - $O(\log(N))$ vs. $O(N)$ in an `unordered_map`
 - Use an `unordered_map` if you want to optimize the average case complexity
 - $O(1)$ vs. $O(\log(N))$ in a `map`

References

- Lippman Chapters 10,11
- <http://www.cplusplus.com/reference/stl/>
- <http://www.learncpp.com/cpp-tutorial>

Credits

- Bjarne Stroustrup. www.stroustrup.com/Programming