

Message Passing Interface

Eugenio Gianniti & Danilo Ardagna

Politecnico di Milano

name.lastname@polimi.it

07/12/2024



**POLITECNICO
DI MILANO**


Content

- Collective Communication
- Reference

COLLECTIVE COMMUNICATION

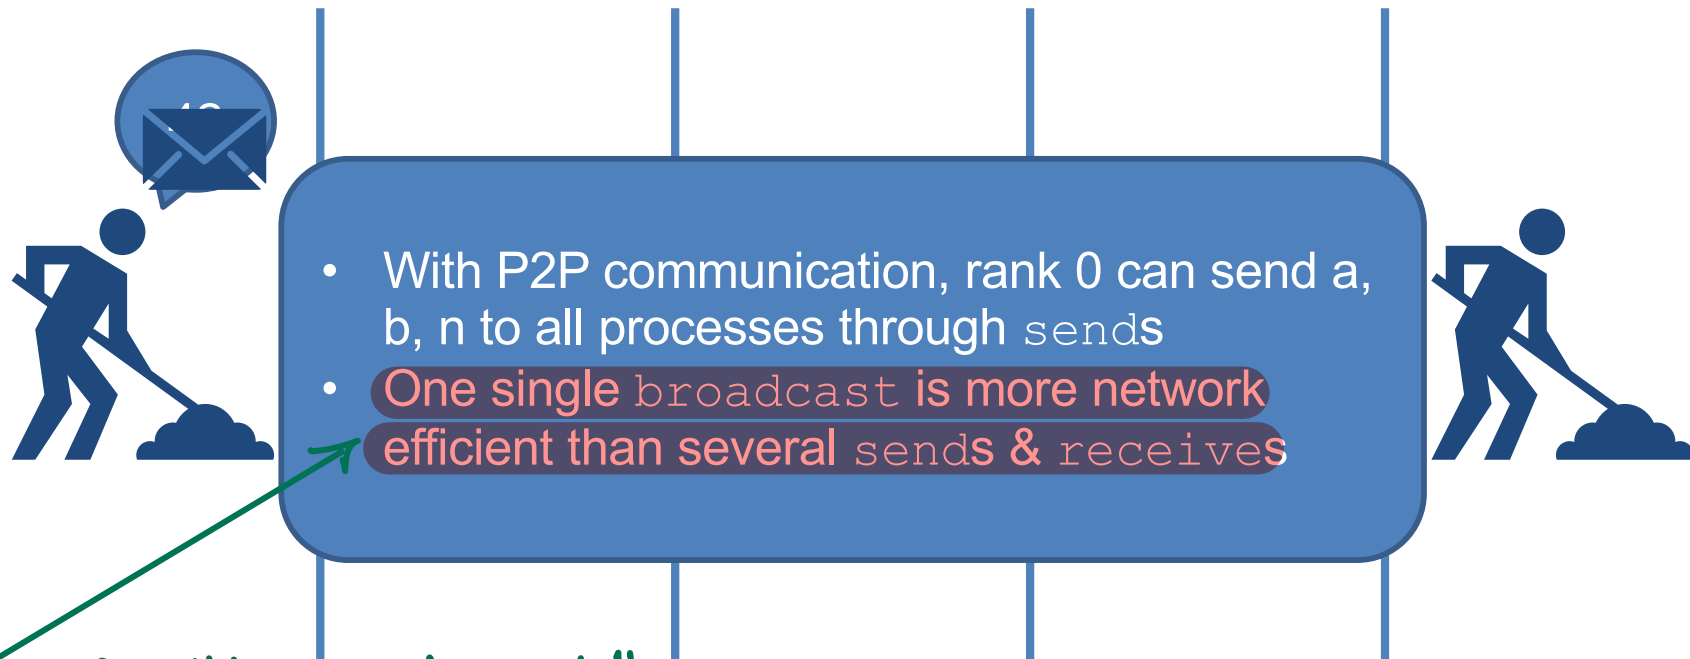
Or **how to throw flyers from a plane**

MPI in Pictures – Send & Receive



MPI Library Runtime

MPI in Pictures – Collective (broadcast)



Interest of "broadcast"



Collective Routines



- Involve all the processes in a communicator
- Only blocking routines



- Transmit only predefined MPI data types
- Cannot use tags to identify messages
- **Attention!** Be sure that every process in the communicator calls the **same collective function** to avoid deadlocks

Broadcast → WHAT DOES IT DO ?

- Delivers an exact copy of the data in buffer from root to all the processes in comm

```
int MPI_Bcast (void *buffer, int count,
               MPI_Datatype datatype,
               int root, MPI_Comm comm)
```

Rank	buffer	
	Before	After
0	A	A
1	?	A
2	?	A
3	?	A

Trapezoidal Rule — Broadcasting Inputs

```
void get_input (double & a, double & b, unsigned & n)
{
    int rank;
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);

    if (rank == 0)
        std::cin >> a >> b >> n;

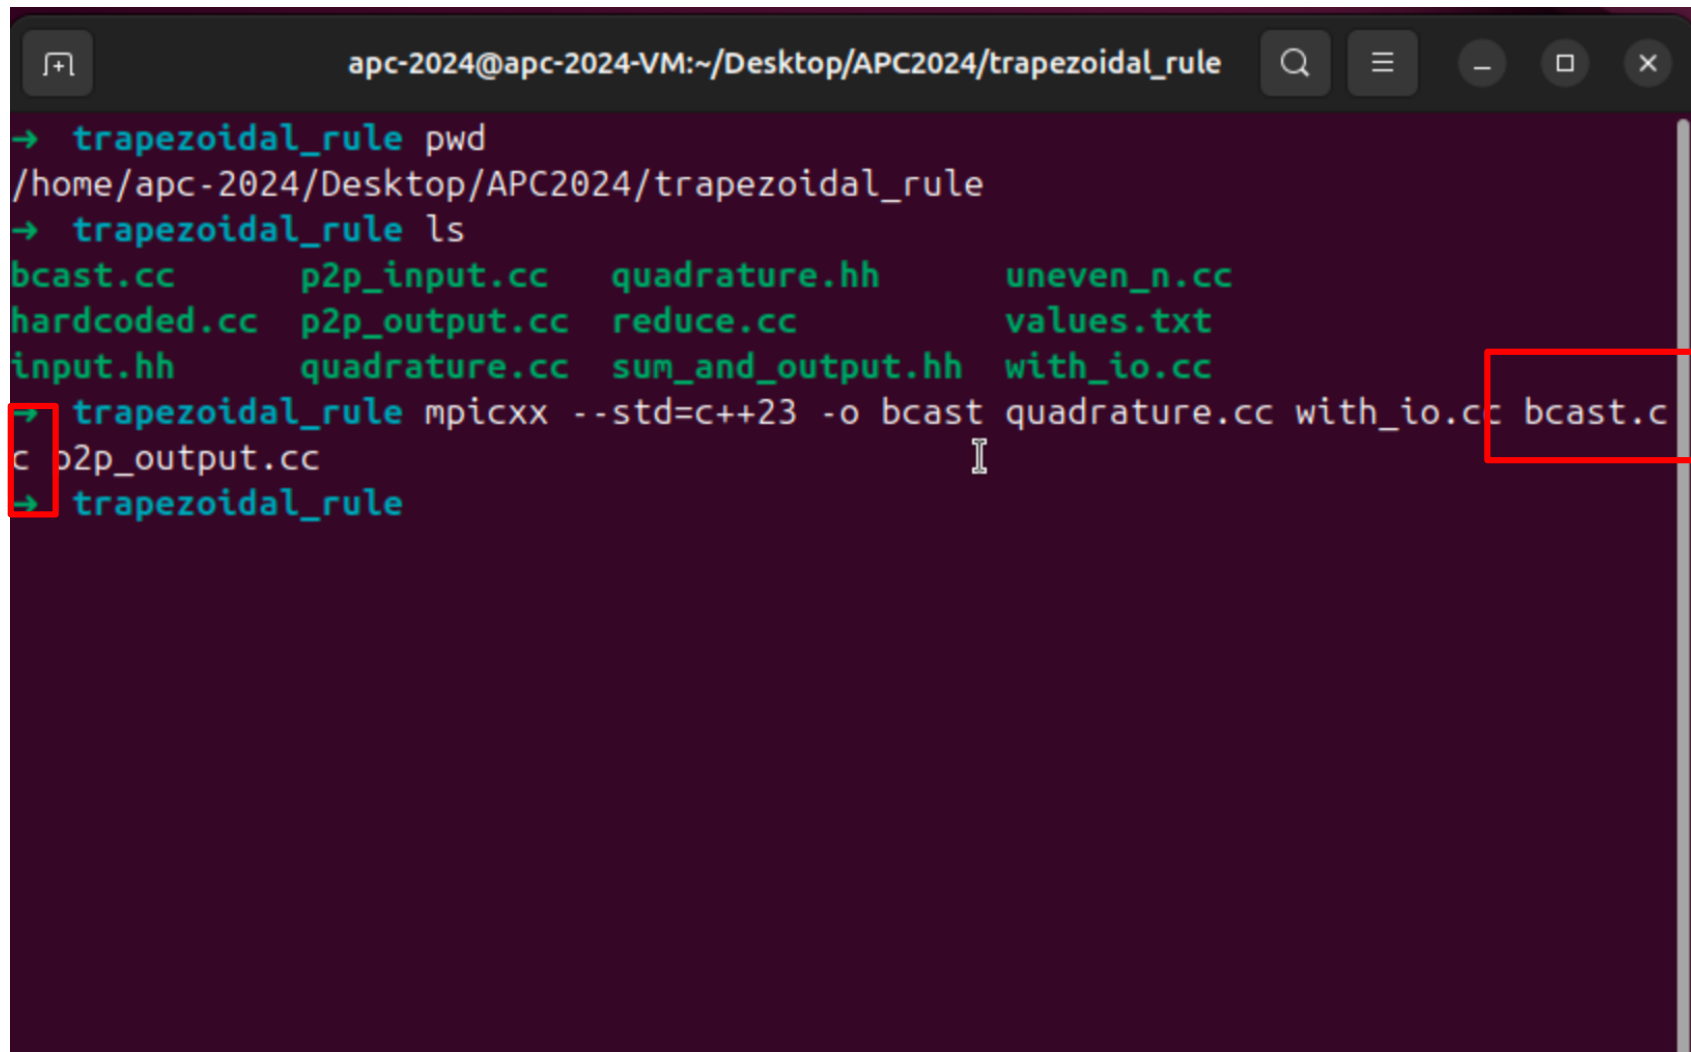
    MPI_Bcast (&a, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);
    MPI_Bcast (&b, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);
    MPI_Bcast (&n, 1, MPI_UNSIGNED, 0, MPI_COMM_WORLD);
}
```



DEMO

Trapezoidal Rule — Broadcasting Inputs

```
void get_input (double & a, double & b, unsigned & n)
```

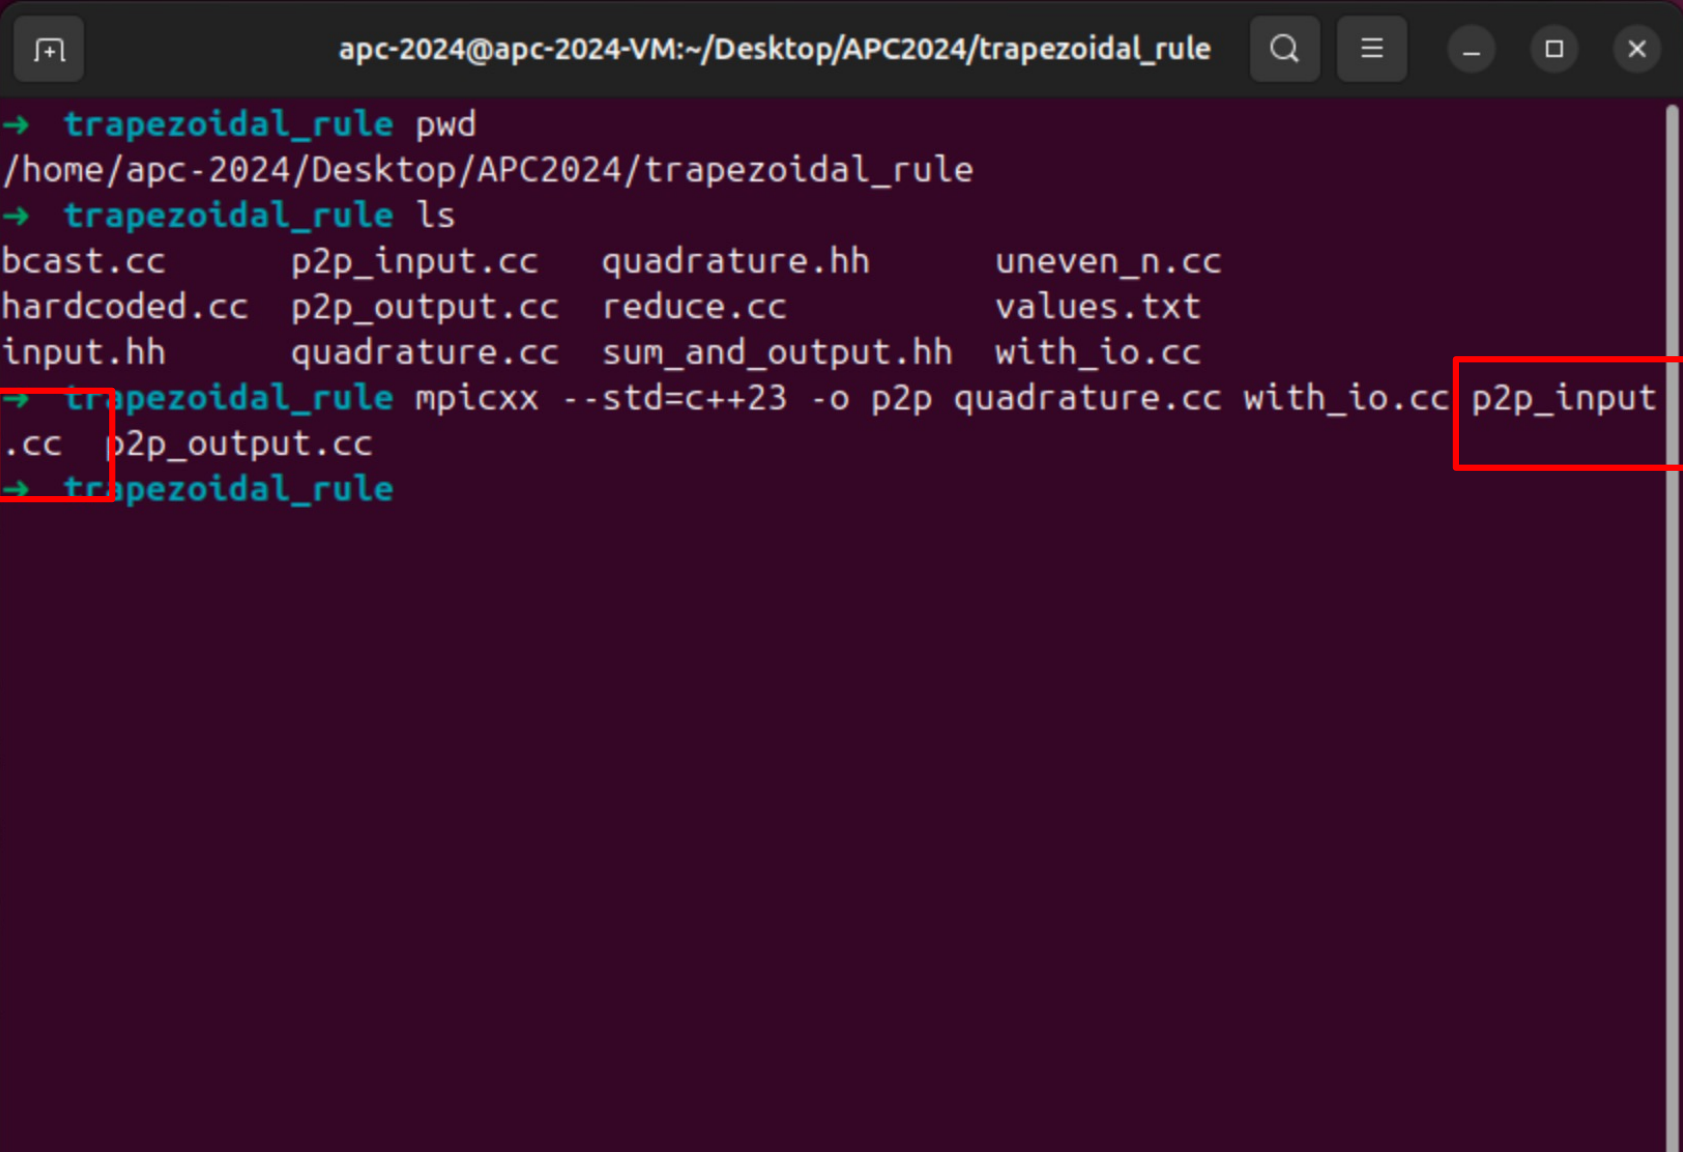


A terminal window titled "apc-2024@apc-2024-VM:~/Desktop/APC2024/trapezoidal_rule" showing the following commands and output:

```
→ trapezoidal_rule pwd
/home/apc-2024/Desktop/APC2024/trapezoidal_rule
→ trapezoidal_rule ls
bcast.cc      p2p_input.cc  quadrature.hh  uneven_n.cc
hardcoded.cc  p2p_output.cc reduce.cc      values.txt
input.hh      quadrature.cc sum_and_output.hh with_io.cc
→ trapezoidal_rule mpicxx --std=c++23 -o bcast quadrature.cc with_io.cc bcast.c
c p2p_output.cc
→ trapezoidal_rule
```

Red boxes highlight the terminal prompt "→" on the first line of the compilation command and the file "bcast.c" in the command arguments.

Trapezoidal Rule — Based on P2P



```
apc-2024@apc-2024-VM:~/Desktop/APC2024/trapezoidal_rule
→ trapezoidal_rule pwd
/home/apc-2024/Desktop/APC2024/trapezoidal_rule
→ trapezoidal_rule ls
bcast.cc      p2p_input.cc  quadrature.hh  uneven_n.cc
hardcoded.cc  p2p_output.cc reduce.cc      values.txt
input.hh      quadrature.cc sum_and_output.hh with_io.cc
→ trapezoidal_rule mpicxx --std=c++23 -o p2p quadrature.cc with_io.cc p2p_input
.cc p2p_output.cc
→ trapezoidal_rule
```

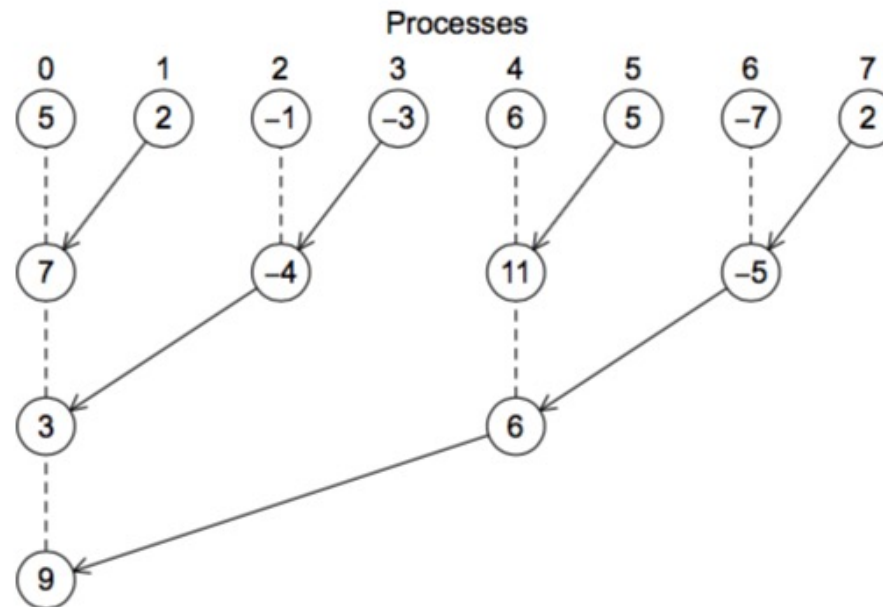
The image shows a terminal window with a dark purple background. The window title is 'apc-2024@apc-2024-VM:~/Desktop/APC2024/trapezoidal_rule'. The user has executed several commands: 'pwd' to show the current directory, 'ls' to list files, and 'mpicxx' to compile the program. The 'ls' command output lists several files, including 'p2p_input.cc' and 'p2p_output.cc'. The 'mpicxx' command is shown with arguments for standard library, output file, and source files. Red boxes highlight the 'p2p_input' part of the command and the 'p2p_output.cc' file in the directory listing.

Trapezoidal Rule - Splitting Sum Work



- We are not computing the “global sum” efficiently
- If we hire N workers to build a house, we might feel that we weren’t getting our money’s worth if $N - 1$ of the workers told the first what to do and then the $N - 1$ collected their pay and went home
 - But this is what we’re doing in our global sum:
 - Each process with rank greater than 0 is “telling process 0 what to do” and then quitting
 - Process 0 is doing nearly all the work in computing the global sum

Trapezoidal Rule - Splitting Sum Work



- This is what is done for us by `MPI_Reduce` (if next slide)
- There are also some communication optimizations behind the curtain, but we don't enter the details

Reduce *SYNTAX*

- Applies `op` to portions of data in `sendbuf` from all the processes in `comm`, storing the result in `recvbuf` on `dest`

```
int MPI_Reduce (const void *sendbuf,  
               void *recvbuf, int count,  
               MPI_Datatype datatype,  
               MPI_Op op, int dest,  
               MPI_Comm comm)
```

cf example @ next slide.

Example of Reduce

```
double local_partial = // some partial sum ;  
double total;  
MPI_Reduce (&local_partial, &total, 1,  
            MPI_DOUBLE, MPI_SUM,  
            0, MPI_COMM_WORLD) ;
```

Rank	local_partial	total
0	1	10
1	2	N/A
2	3	N/A
3	4	N/A

Reduce Operators

MPI_Op


- The standard provides several ready to use operators for reduce routines

MPI_MAX	MPI_LAND	MPI_LXOR
MPI_MIN	MPI_BAND	MPI_BXOR
MPI_SUM	MPI_LOR	MPI_MAXLOC
MPI_PROD	MPI_BOR	MPI_MINLOC

We are interested in only on these ones

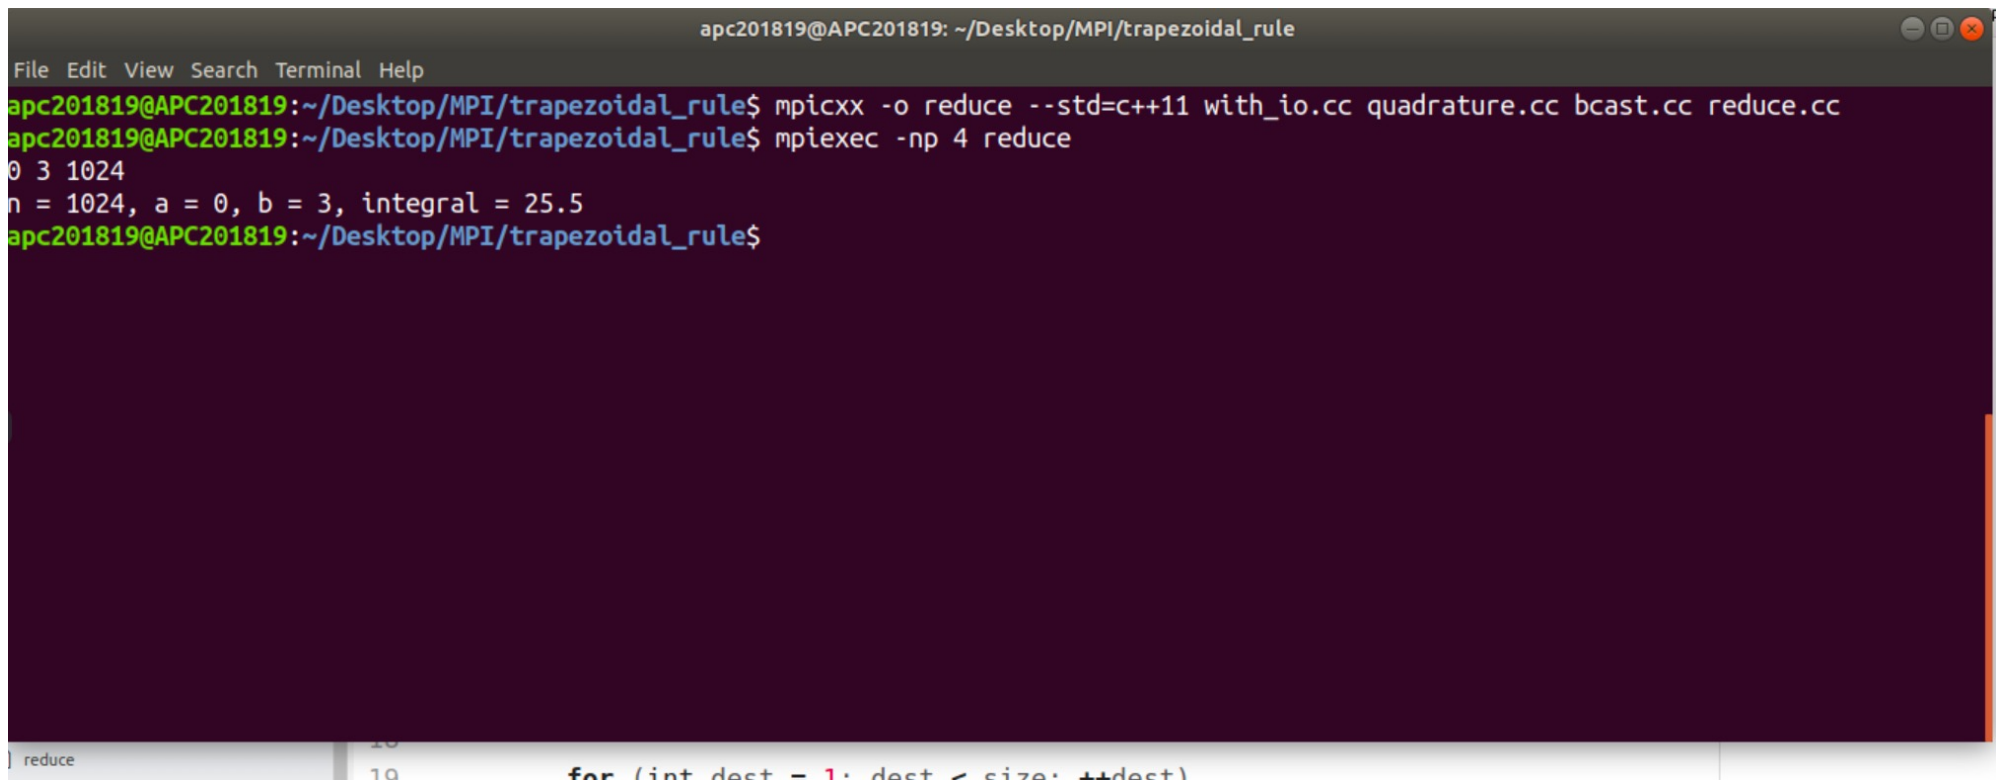
Trapezoidal Rule - Splitting Sum Work

```
void sum_and_print (double local_integral,  
                    std::ostream & out,  
                    double a, double b, unsigned n)  
{  
    int rank;  
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);  
    double total (0.);  
    MPI_Reduce (&local_integral, &total, 1, MPI_DOUBLE,  
               MPI_SUM, 0, MPI_COMM_WORLD);  
    if (rank == 0)  
    {  
        out << "n = " << n  
              << ", a = " << a << ", b = " << b  
              << ", integral = " << total << std::endl;  
    }  
}
```



Trapezoidal Rule - Splitting Sum Work

```
void sum_and_print (double local_integral,  
                   std::ostream & out,  
                   double a, double b, unsigned n)  
{
```



The screenshot shows a terminal window titled "apc201819@APC201819: ~/Desktop/MPI/trapezoidal_rule". The terminal contains the following commands and output:

```
File Edit View Search Terminal Help
apc201819@APC201819:~/Desktop/MPI/trapezoidal_rule$ mpicxx -o reduce --std=c++11 with_io.cc quadrature.cc bcast.cc reduce.cc
apc201819@APC201819:~/Desktop/MPI/trapezoidal_rule$ mpiexec -np 4 reduce
0 3 1024
n = 1024, a = 0, b = 3, integral = 25.5
apc201819@APC201819:~/Desktop/MPI/trapezoidal_rule$
```

At the bottom of the terminal window, a snippet of C++ code is visible:

```
for (int dest = 1; dest < size; ++dest)
```

More on MPI_Reduce

about "count" • By using a count argument greater than 1, MPI_Reduce can operate on arrays instead of scalars

- The following code could thus be used to add a collection of N -dimensional vectors:

```
std::vector<double> local_x (N), sum (N);
/* partial computation on local_x */
MPI_Reduce (local_x.data (), sum.data (),
            N, MPI_DOUBLE, MPI_SUM,
            0, MPI_COMM_WORLD);
```

Returns the pointer to the vector elements

MPI_Allreduce

context
↓
new
context

- In our trapezoidal rule program, we just print the result, so it's perfectly natural for only one process to get the result of the global sum
- In some situations (e.g., you are implementing a **parallel function**) **all of the processes** might need the result of a global sum in order to complete some larger computation
- MPI provides a variant of `MPI_Reduce` that **will store the result on all the processes in the communicator**

```
int MPI_Allreduce (const void *sendbuf,  
                  void *recvbuf, int count,  
                  MPI_Datatype datatype,  
                  MPI_Op op, MPI_Comm comm)
```

Collective Communication “In Place”

- Collective communication routines generally use both a send and a receive buffer
 - When dealing with lots of data this implies that you are also occupying a lot of memory
- MPI provides the special placeholder `MPI_IN_PLACE` to enable the use of a single buffer for both input and output


```
double minimum (local_min);  
MPI_Allreduce (MPI_IN_PLACE, &minimum, 1,  
               MPI_DOUBLE, MPI_MIN,  
               MPI_COMM_WORLD);
```

↙ only one buffer.

Collective Communication “In Place”

- It might be tempting to call `MPI_Reduce` using the same buffer for both input and output

```
MPI_Reduce(&x , &x, 1, MPI_DOUBLE, MPI_SUM, 0, comm);
```



- This call is illegal in MPI (undefined behavior), its result will be unpredictable: it might produce an incorrect result, it might cause the program to crash, it might even produce a correct result

Instead, use `MPI_Allreduce` (cf previous slide).

- It's illegal because it involves **aliasing** of an output argument
 - Two arguments are aliased if they refer to the same block of memory
 - MPI prohibits aliasing of arguments if one of them is an output or input/output argument

*and
next slide.*

Collective Communication “In Place”

- It might be tempting to call `MPI_Reduce` using the same buffer for both input and output

```
MPI_Reduce(&x , &x, 1, MPI_DOUBLE, MPI_SUM, 0, comm);
```

Use `MPI_IN_PLACE` instead:

```
MPI_Reduce(MPI_IN_PLACE, &x, 1, MPI_DOUBLE, MPI_SUM,  
0, comm);
```

- It's illegal because it involves **aliasing** of an output argument
 - Two arguments are aliased if they refer to the same block of memory
 - MPI prohibits aliasing of arguments if one of them is an output or input/output argument

Working With Vectors

Scatter & Gather

Data Distributions

- Suppose we want to write a function that computes a vector sum:

$$\begin{aligned}\mathbf{x} + \mathbf{y} &= (x_0, x_1, \dots, x_{n-1}) + (y_0, y_1, \dots, y_{n-1}) \\ &= (x_0 + y_0, x_1 + y_1, \dots, x_{n-1} + y_{n-1}) = (z_0, z_1, \dots, z_{n-1}) = \mathbf{z}\end{aligned}$$

```
// Assumption: x and y with the same size
vector<double> sum (const vector<double> & x,
                   const vector<double> & y)
{
    vector<double> z(x.size ());
    for (std::size_t i = 0; i < x.size (); ++i)
        z[i] = x[i] + y[i];
    return z;
}
```

→ classical way to do it (no MPI).

Data Distributions

- If the number of components is n and we have comm_sz cores or processes, let's assume that n evenly divides comm_sz and define $\text{local_n} = n / \text{comm_sz}$
- We can parallelize the sum by assigning blocks of local_n consecutive components to each process

Process	Block			
0	0	1	2	3
1	4	5	6	7
2	8	9	10	11

- This is the so-called **block partition**


Data Distributions

- An alternative to a block partition is a **cyclic partition**, where components are assigned in a round robin fashion

Process	Cyclic			
0	0	3	6	9
1	1	4	7	10
2	2	5	8	11

```
for (size_t i=rank; i< v.size(); i+=size) {  
    // do something on v[i]  
}
```

→ Here: "+3"



Data Distributions

- Back to the vector sum problem, given the partitioning scheme, each process simply adds its assigned components
- **Block partitioning** is used when **data source** is available on a **single process**
 - We have a primitive for this (`MPI_Scatter`)
- **Cyclic partitioning** is used when data source is already available across all processes
 - This is general, we do not need even any assumptions on the number of elements to be processed

Data Distributions - Block Partitioning

↳ 1st method : cf slide 25

- Each process will have `local_n` components of the vectors, and, in order to save on storage, we can just store these on each process as a vector of `local_n` elements

```
vector<double> parallel_sum (const vector<double> &
local_x, const vector<double> & local_y)
{
    vector<double> local_z(local_x.size ());
    for (size_t i = 0; i < local_x.size (); ++i)
        local_z[i] = local_x[i] + local_y[i];
    return local_z;
}
```

of size `local_n` { we store only `local_x`
& `local_y` on each process.

Scatter - Block Partitioning

- To implement our vector addition function, we need to:

- read the dimension of the vectors
- then read in the vectors **x** and **y**

- Process 0 can prompt the user, read in the dimension, and broadcast the value to the other processes

- For the vectors:

- process 0 reads them (no other option)
- then process 0 sends the needed components to each of the other processes

- This is exactly what MPI Scatter implements (block partition scheme)

↑ MPI_Scatter = block partition implementation ↑

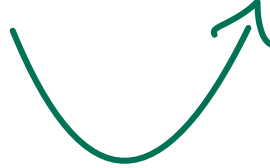
Scatter WHAT DOES IT DO ?

- Sends a portion of the data in `sendbuf` from `root` to all the processes in `comm`, storing it in `recvbuf`
- `sendbuf` holds `sendcount * size` elements (in other words, `sendcount` is the number of elements sent to individual processes)

```
int MPI_Scatter (const void *sendbuf,  
                int sendcount,  
                MPI_Datatype sendtype,  
                void *recvbuf,  
                int recvcount,  
                MPI_Datatype recvtype,  
                int root, MPI_Comm comm)
```

Example of Scatter – rank 0 is root

Rank	sendbuf	recvbuf
0	ABCDEFGH	AB
1	N/A	CD
2	N/A	EF
3	N/A	GH



clearly it's the
implementation
of block partition.

read_vector ^{1/3}

```
std::vector<double>
read_vector (unsigned n,
             std::string const & name,
             MPI_Comm const & comm)
{
    int rank, size;
    MPI_Comm_rank (comm, &rank);
    MPI_Comm_size (comm, &size);
    const unsigned local_n = n / size;
    std::vector<double> result (local_n);
```

Extremely
important to
place this here

read_vector 2/3

```
if (rank == 0)
{
    std::vector<double> input (n);
    std::cout << "Enter " << name << "\n";
    for (double & e : input)
        std::cin >> e;
    MPI_Scatter (input.data (), local_n,
                MPI_DOUBLE,
                result.data (), local_n,
                MPI_DOUBLE, 0, comm);
}
```

`read_vector`

3/3

```
else
```

```
{
```

```
    /* Here are only receiving ranks,  
     * no need for the send buffer. */
```

```
    MPI_Scatter (nullptr, local_n,  
                MPI_DOUBLE,  
                result.data (), local_n,  
                MPI_DOUBLE, 0, comm);
```

```
}
```

```
return result;
```

```
}
```

Gather WHAT IS IT ?

- Joins portions of data in `sendbuf` from all the processes in `comm` to `root`, storing them all in `recvbuf`
- `recvcount` values received from each process

```
int MPI_Gather (const void *sendbuf,
```

- `MPIAll_gather` provides the destination buffer to all processes
- Same parameters as `MPI_Gather` only `root` is missing

Gather WHAT IS IT ?

- Joins portions of data in `sendbuf` from all the processes in `comm` to `root`, storing them all in `recvbuf`
- `recvcount` values received from each process

```
int MPI_Gather (const void *sendbuf,
               int sendcount,
               MPI_Datatype sendtype,
               void *recvbuf,
               int recvcount,
               MPI_Datatype recvtype,
               int root, MPI_Comm comm)
```

cf next slide for example.

Example of Gather – rank 0 is root

Rank	sendbuf	recvbuf
0	AB	ABCDEFGH
1	CD	N/A
2	EF	N/A
3	GH	N/A

print_vector 1/3

```
void  
print_vector (vector<double> const &  
              local_v, unsigned n,  
              string const & title,  
              MPI_Comm const & comm)  
{  
    int rank, size;  
    MPI_Comm_rank (comm, &rank);  
    MPI_Comm_size (comm, &size);  
    const unsigned local_n = local_v.size ();
```

print_vector ^{2/3}

```
if (rank > 0)
{
    /* Here are only sending ranks,
     * no need for the receive buffer. */
    MPI_Gather (local_v.data (), local_n,
               MPI_DOUBLE,
               nullptr, local_n,
               MPI_DOUBLE, 0, comm);
}
```

print_vector

```
else
{
```

// Rank 0 receives everything

```
    std::vector<double> global (n);
    MPI_Gather (local_v.data (), local_n,
               MPI_DOUBLE,
               global.data (), local_n,
               MPI_DOUBLE, 0, comm);
```

```
    std::cout << title << "\n";
    for (double value : global)
        std::cout << value << " ";
    std::cout << std::endl;
}
```

*// Reads all the
// values of the vector.*

Final Remarks on Collective Communications



- All the processes in the communicator must call the same collective function
 - If a program attempts to match a call to `MPI_Reduce` on one process with a call to `MPI_Recv` on another process it is erroneous and probably will hang or crash.
- The arguments passed by each process to an MPI collective communication must be “compatible”
 - If one process passes in 0 as the dest process and another passes in 1, then the outcome of a call to, e.g., `MPI_Reduce` is erroneous and the program is likely to hang or crash
- The `recvbuf` argument is only used on `dest` process. However, all of the processes still need to pass in an actual argument corresponding to `recvbuf`, even if it's just `nullptr`

Final Remarks on Collective Communications

- Point-to-point communications are matched based on tags and communicators. Collective communications don't use tags, so they're matched solely based on the communicator and the order in which they're called

Time	Process 0	Process 1	Process 2
0	<code>a = 1; c = 2;</code>	<code>a = 1; c = 2;</code>	<code>a = 1; c = 2;</code>
1	<code>MPI_Reduce(&a,&b,..., 0,comm);</code>	<code>MPI_Reduce(&c,&d,..., 0,comm);</code>	<code>MPI_Reduce(&a,&b,..., 0,comm);</code>
2	<code>MPI_Reduce(&c,&d,..., 0,comm);</code>	<code>MPI_Reduce(&a,&b,..., 0,comm)</code>	<code>MPI_Reduce(&c,&d,..., 0,comm)</code>

Final Remarks on Collective Communications

- Point-to-point communications are matched based on tags and communicators. Collective communications don't use tags, so they're matched solely based on the communicator and the order in which they're called

Time	Process 0	Process 1	Process 2
0	<code>a = 1; c = 2;</code>	<code>a = 1; c = 2;</code>	<code>a = 1; c = 2;</code>
1	<code>MPI_Reduce(&a,&b,..., 0,comm);</code>	<code>MPI_Reduce(&c,&d,..., 0,comm);</code>	<code>MPI_Reduce(&a,&b,..., 0,comm);</code>
2	<code>MPI_Reduce(&c,&d,..., 0,comm);</code>	<code>MPI_Reduce(&a,&b,..., 0,comm)</code>	<code>MPI_Reduce(&c,&d,..., 0,comm)</code>

$$b = 1 + 2 + 1 = 4$$

$$d = 2 + 1 + 2 = 5$$

Reference

- The Open MPI documentation: <https://www.open-mpi.org/doc/current/>
- The MPI tutorial by the Lawrence Livermore National Laboratory: <https://computing.llnl.gov/tutorials/mpi/>
- Pacheco Chapter 3